



Data Science Meetup Kassel
Convolutional Neural Networks
27.02.2018

Nicolas Kuhaupt
Data Scientist Fraunhofer IEE

1. Neuronen

2. Neuronale Netzwerke

3. Convolutional Neural Network (CNN)

4. Anwendungen von CNN

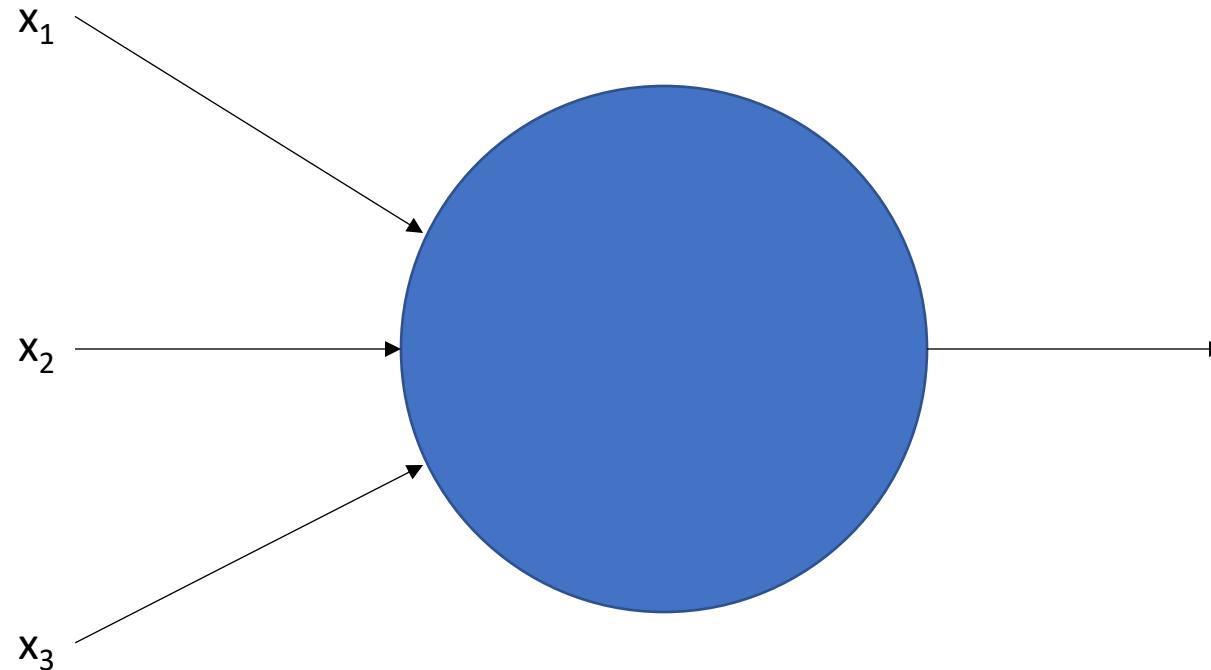
1. Neuronen

2. Neuronale Netzwerke

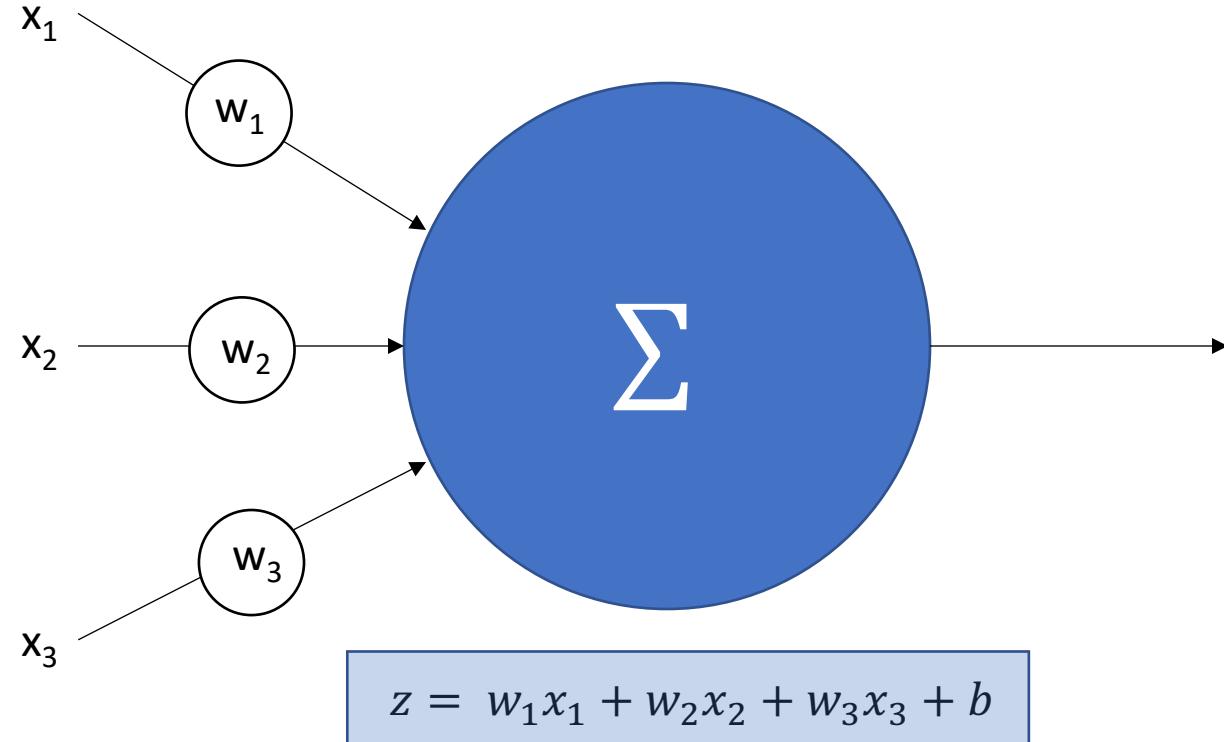
3. Convolutional Neural Network (CNN)

4. Anwendungen von CNN

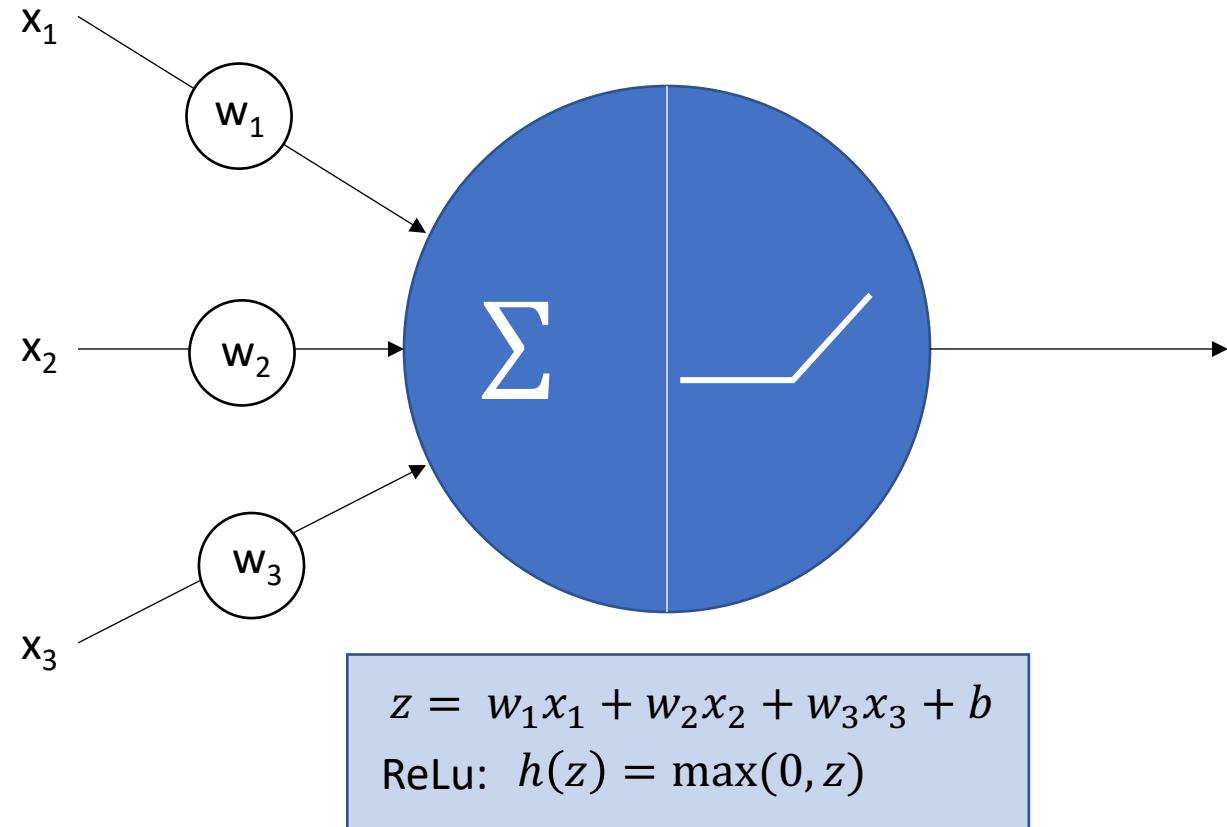
Neuron



Neuron



Neuron - Aktivierungsfunktion



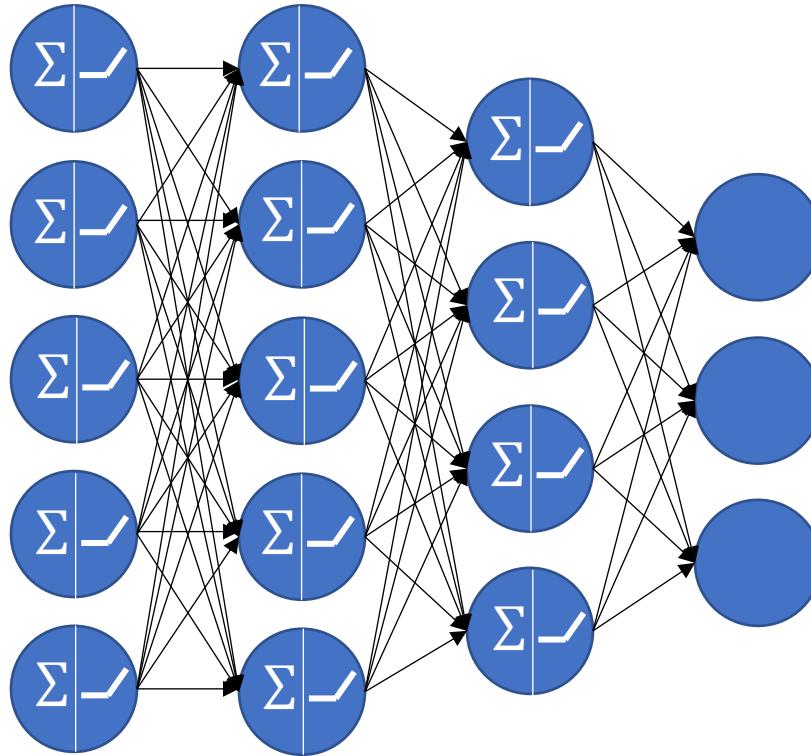
1. Neuronen

2. Neuronale Netzwerke

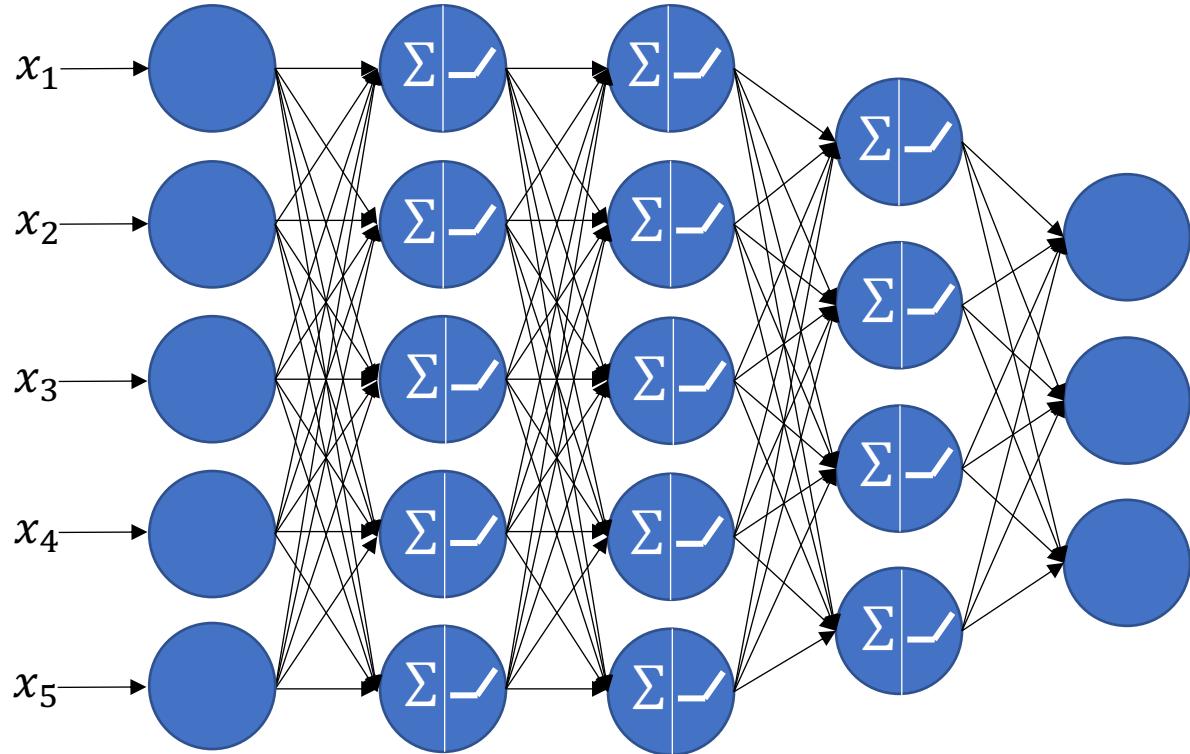
3. Convolutional Neural Network (CNN)

4. Anwendungen von CNN

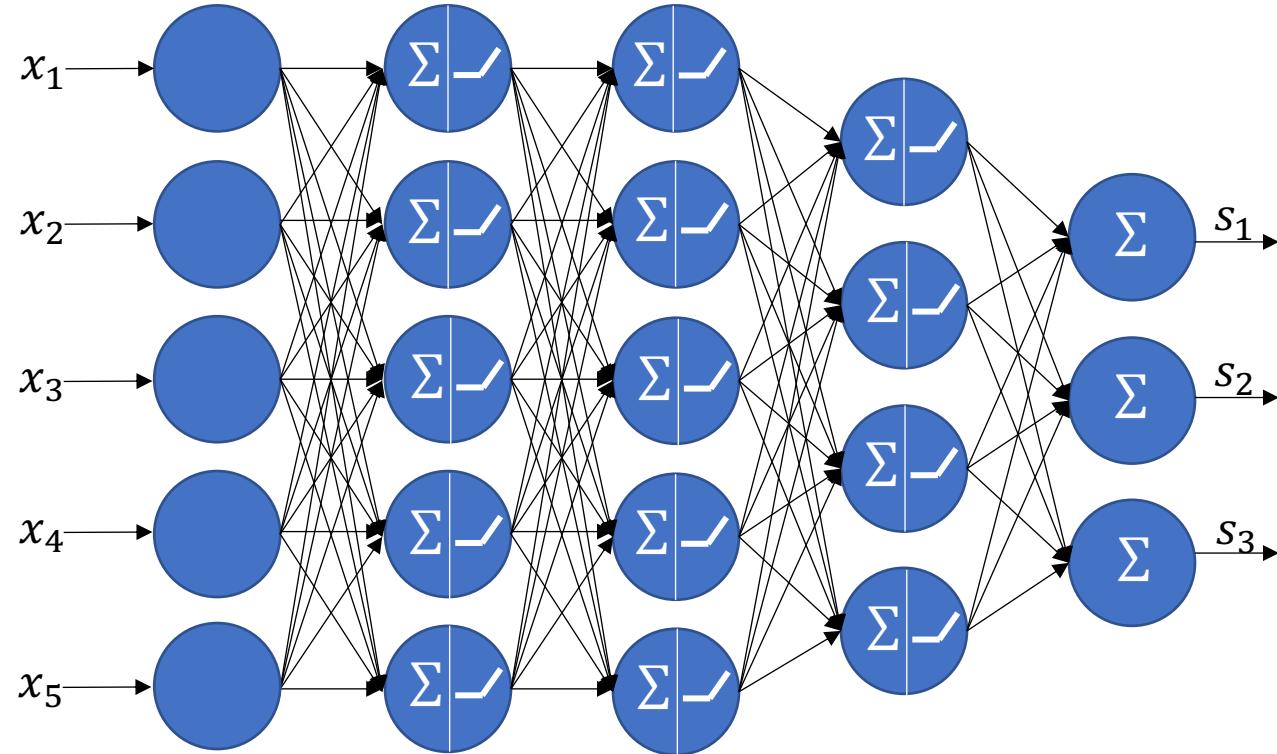
Neuronale Netze



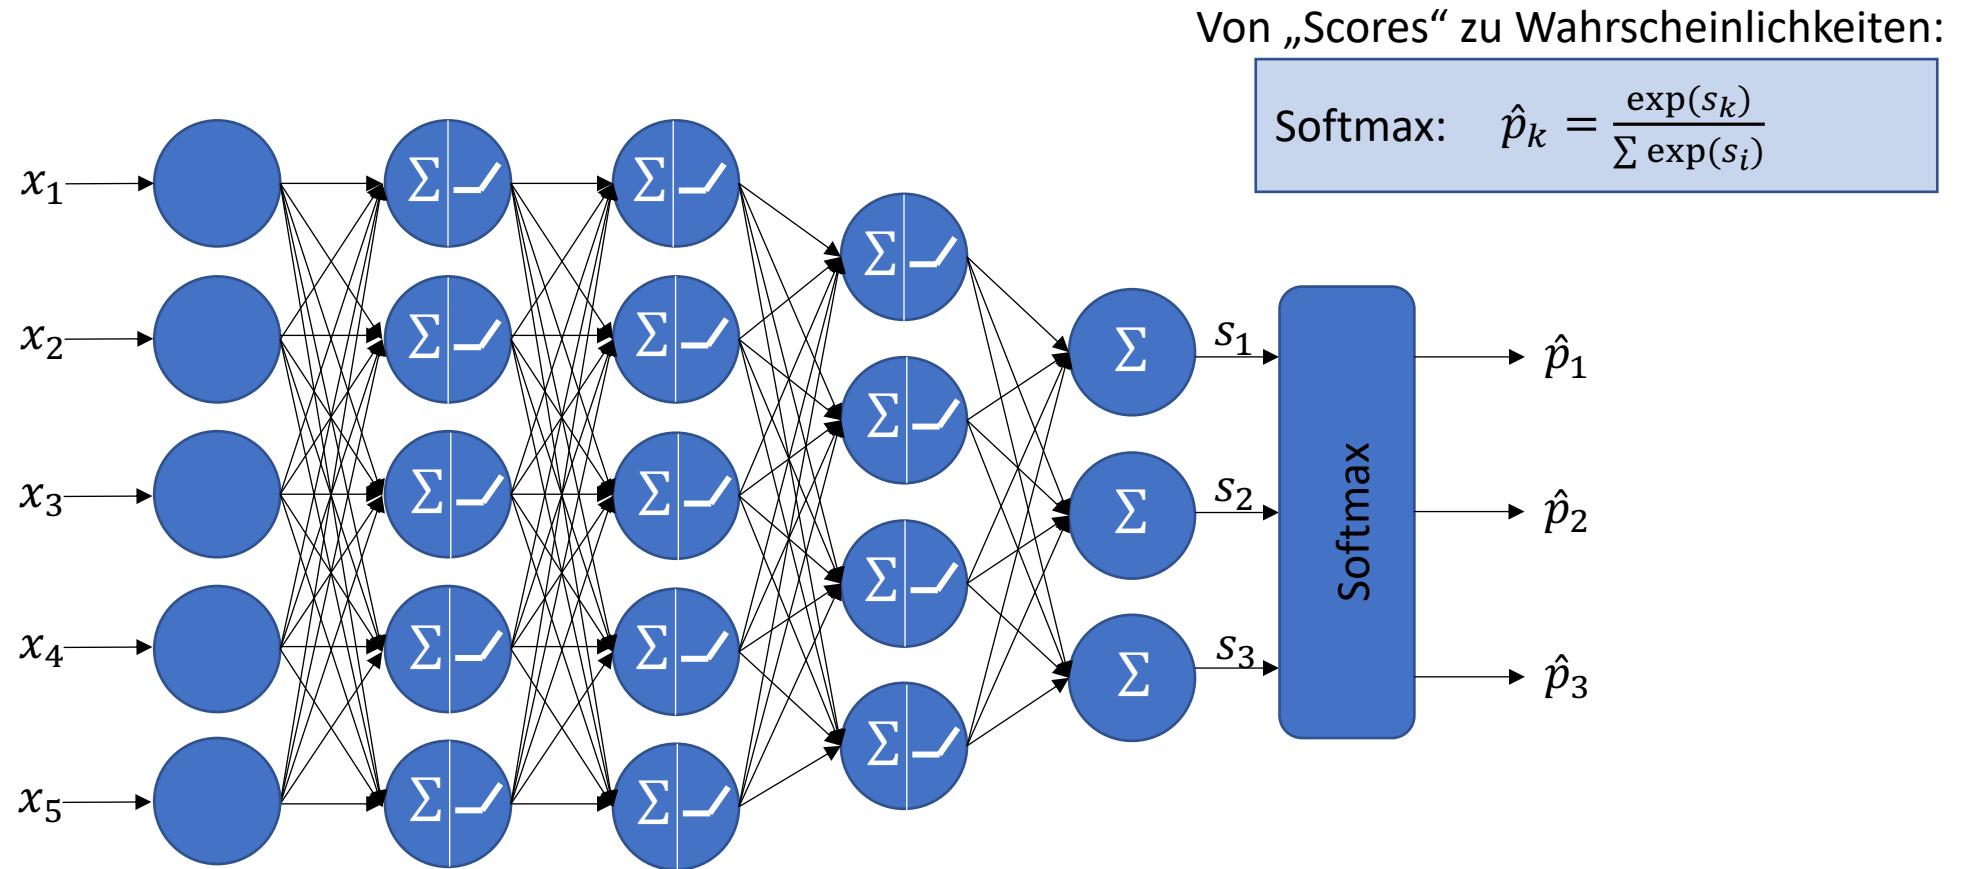
Neuronale Netze – Input Layer



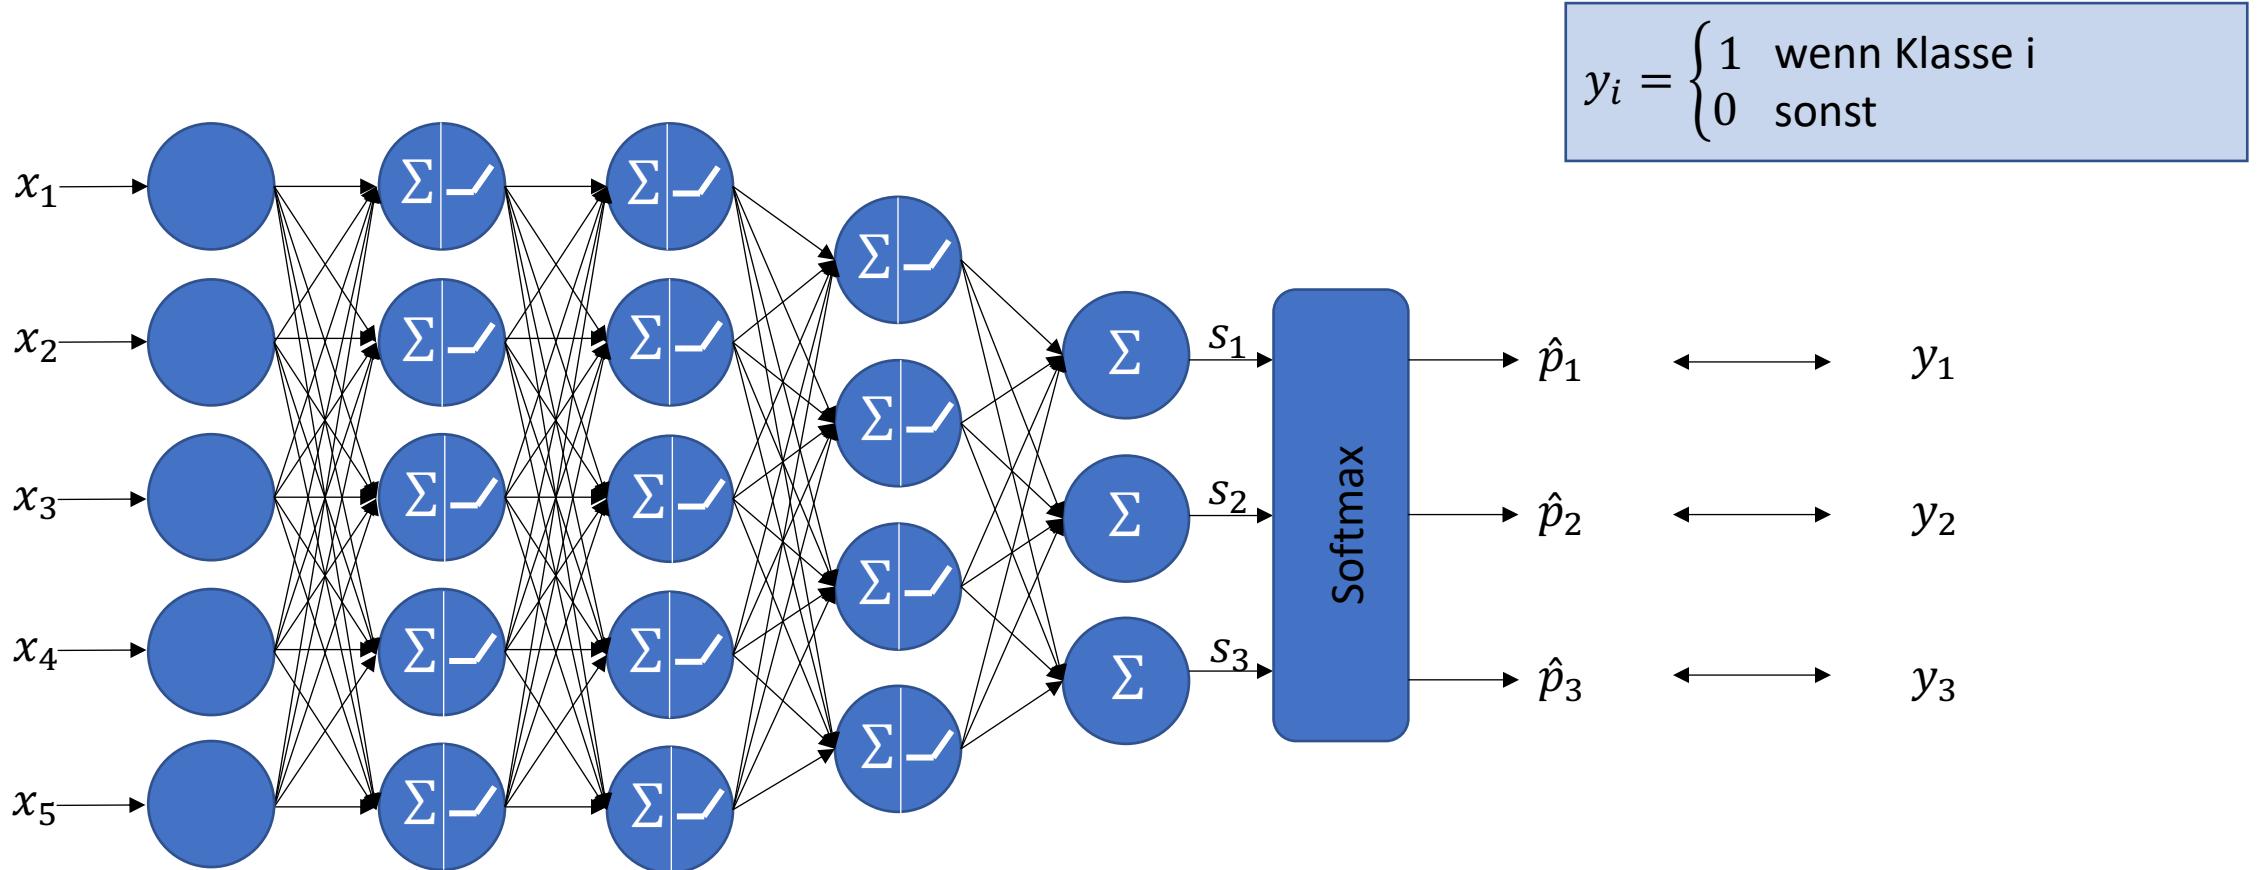
Neuronale Netze – Output Layer (I)



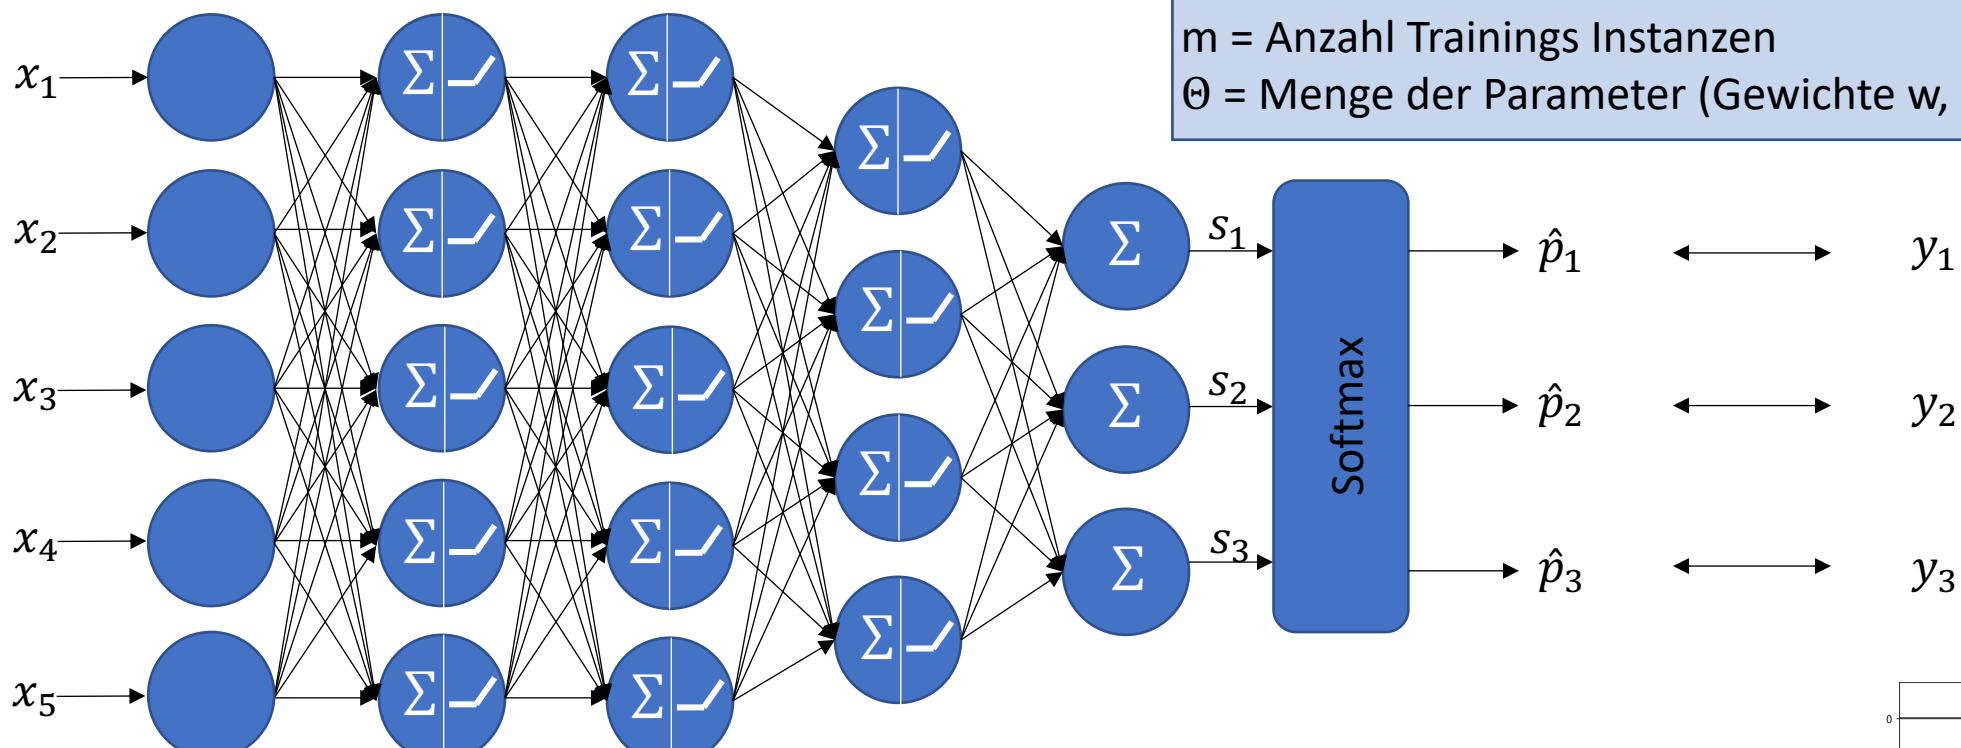
Neuronale Netze – Output Layer (II)



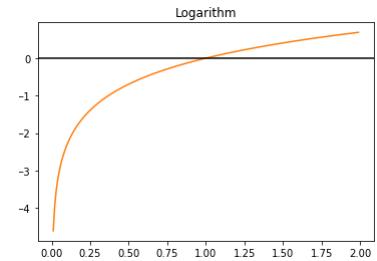
Neuronale Netze – Wie lernt ein NN?



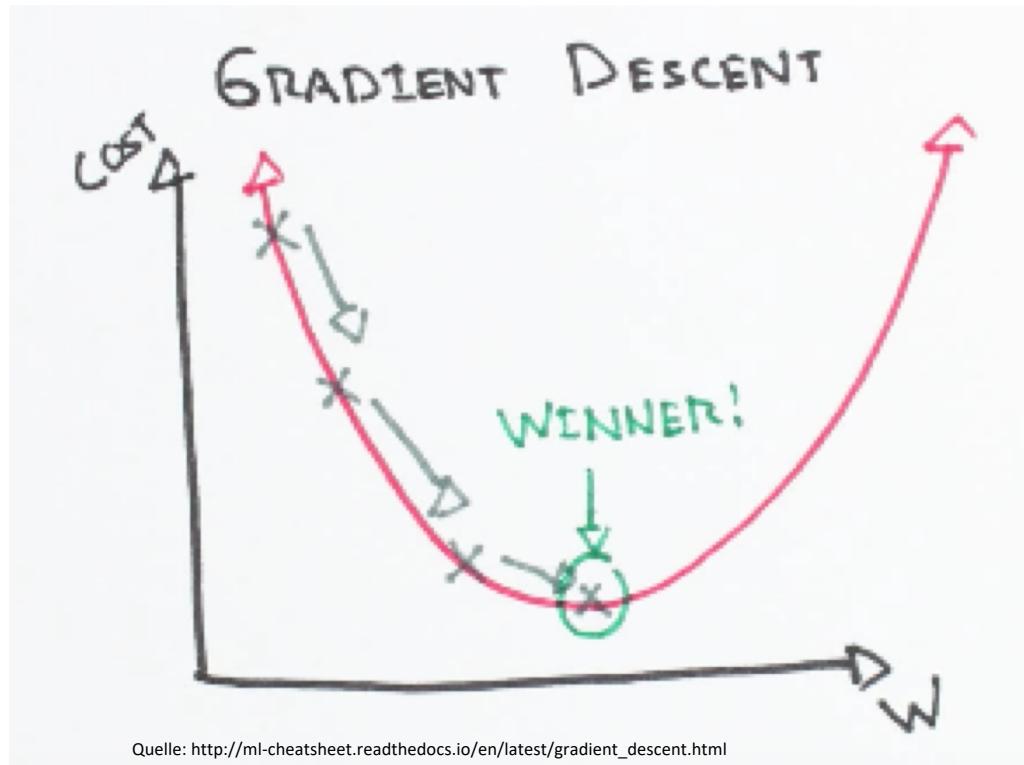
Neuronale Netze – Kostenfunktion



Cross Entropy: $C(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(\hat{p}_k^{(i)})$
K = Anzahl Klassen
m = Anzahl Trainings Instanzen
 Θ = Menge der Parameter (Gewichte w, Bias b)



Neuronale Netze – Gradient Descent (I)



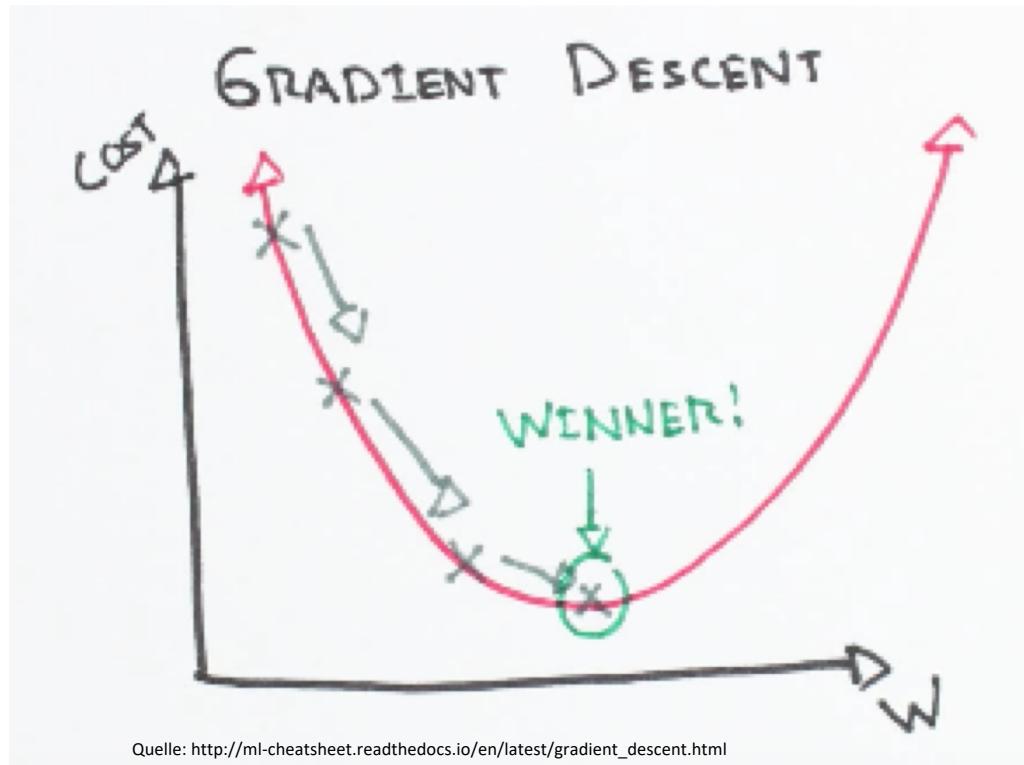
Cross Entropy: $C(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(\hat{p}_k^{(i)})$

K = Anzahl Klassen

m = Anzahl Trainings Instanzen

Θ = Menge der Parameter (Gewichte w, Bias b)

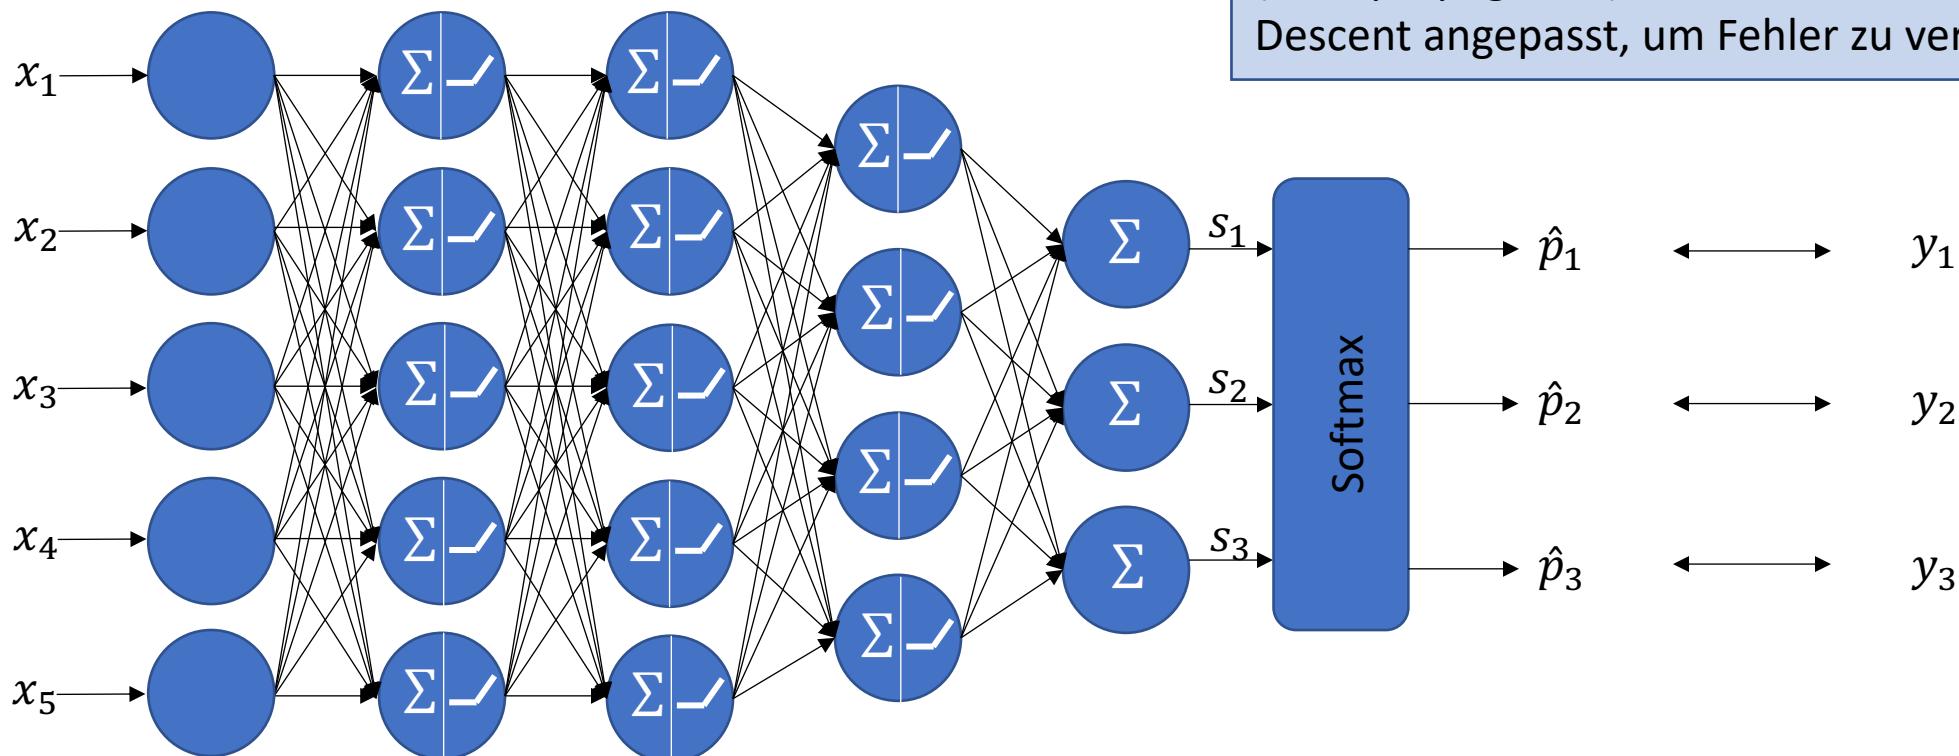
Neuronale Netze – Gradient Descent (II)



Cross Entropy: $C(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(\hat{p}_k^{(i)})$
K = Anzahl Klassen
m = Anzahl Trainings Instanzen
 Θ = Menge der Parameter (Gewichte w, Bias b)

Stochastic Gradient Descent: m = 1
Mini Batch Gradient Descent: z.B. m = 50
Batch Gradient Descent: m = n

Neuronale Netze – Backpropagation



Tensorflow Intermezzo



- Deep Learning API in Python
- In C++ implementiert
- Multiplattform
- Visualisierung (Tensorboard)
- High Level API on top (Keras)
- Graph und Session

MNIST Datensatz

- Zahlen 0 – 9
- 28 x 28 Pixel
- 60000 Daten
- Jeder Pixel entspricht einem Grau-Wert



Quelle: https://en.wikipedia.org/wiki/MNIST_database#/media/File:MnistExamples.png

Tensorflow – NN Implementation: Graph

```
import tensorflow as tf

n_inputs = 28*28
n_hidden1 = 300
n_hidden2 = 100
n_outputs = 10

x = tf.placeholder(tf.float32, shape=(None, n_inputs), name="X")
y = tf.placeholder(tf.int64, shape=(None), name="y")

with tf.name_scope("dnn"):
    hidden1 = tf.layers.dense(x, n_hidden1, name="hidden1", activation=tf.nn.relu)
    hidden2 = tf.layers.dense(hidden1, n_hidden2, name="hidden2", activation=tf.nn.relu)
    logits = tf.layers.dense(hidden2, n_outputs, name="outputs")

with tf.name_scope("loss"):
    xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y, logits=logits, name="xentropy")
    loss = tf.reduce_mean(xentropy, name="loss")

learning_rate = 0.01

with tf.name_scope("train"):
    optimizer = tf.train.GradientDescentOptimizer(learning_rate)
    training_op = optimizer.minimize(loss)

with tf.name_scope("eval"):
    correct = tf.nn.in_top_k(logits, y, 1)
    accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
```

● ● ● ● ... ● ● ● ● ● 784

● ● ● ● ... ● ● ● ● ● 300
● ● ● ● ... ● ● ● ● ● 200
● ● ● ● 10

y_0, y_1, \dots, y_9

Tensorflow – NN Implementation: Session

```
from tensorflow.examples.tutorials.mnist import input_data

mnist = input_data.read_data_sets("/tmp/data/")

init = tf.global_variables_initializer()

n_epochs = 20
batch_size = 50

with tf.Session() as sess:
    init.run()
    for epoch in range(n_epochs):
        for iteration in range(mnist.train.num_examples // batch_size):
            X_batch, y_batch = mnist.train.next_batch(batch_size)
            sess.run(training_op, feed_dict={X: X_batch, y: y_batch})
            acc_train = accuracy.eval(feed_dict={X: X_batch, y: y_batch})
            acc_test = accuracy.eval(feed_dict={X: mnist.test.images, y: mnist.test.labels})
            print(epoch, "Train accuracy:", acc_train, "Test accuracy:", acc_test)

0 Train accuracy: 0.92 Test accuracy: 0.9061
1 Train accuracy: 0.98 Test accuracy: 0.9219
2 Train accuracy: 0.94 Test accuracy: 0.9326
3 Train accuracy: 0.82 Test accuracy: 0.938
4 Train accuracy: 0.9 Test accuracy: 0.9432
5 Train accuracy: 0.96 Test accuracy: 0.9468
6 Train accuracy: 0.98 Test accuracy: 0.9503
7 Train accuracy: 0.94 Test accuracy: 0.9539
8 Train accuracy: 0.98 Test accuracy: 0.9569
9 Train accuracy: 1.0 Test accuracy: 0.9584
10 Train accuracy: 0.98 Test accuracy: 0.9592
```

1. Neuronen

2. Neuronale Netzwerke

3. Convolutional Neural Network (CNN)

4. Anwendungen von CNN

Convolutional Neural Network – Pooling Layer

1	3	4	4
7	8	3	5
7	0	6	2
2	4	6	9



8	5
7	9

Convolutional Neural Network – Pooling Layer

1	3	4	4
7	8	3	5
7	0	6	2
2	4	6	9



8	5
7	9

Convolutional Neural Network – Pooling Layer

1	3	4	4
7	8	3	5
7	0	6	2
2	4	6	9



8	5
7	9

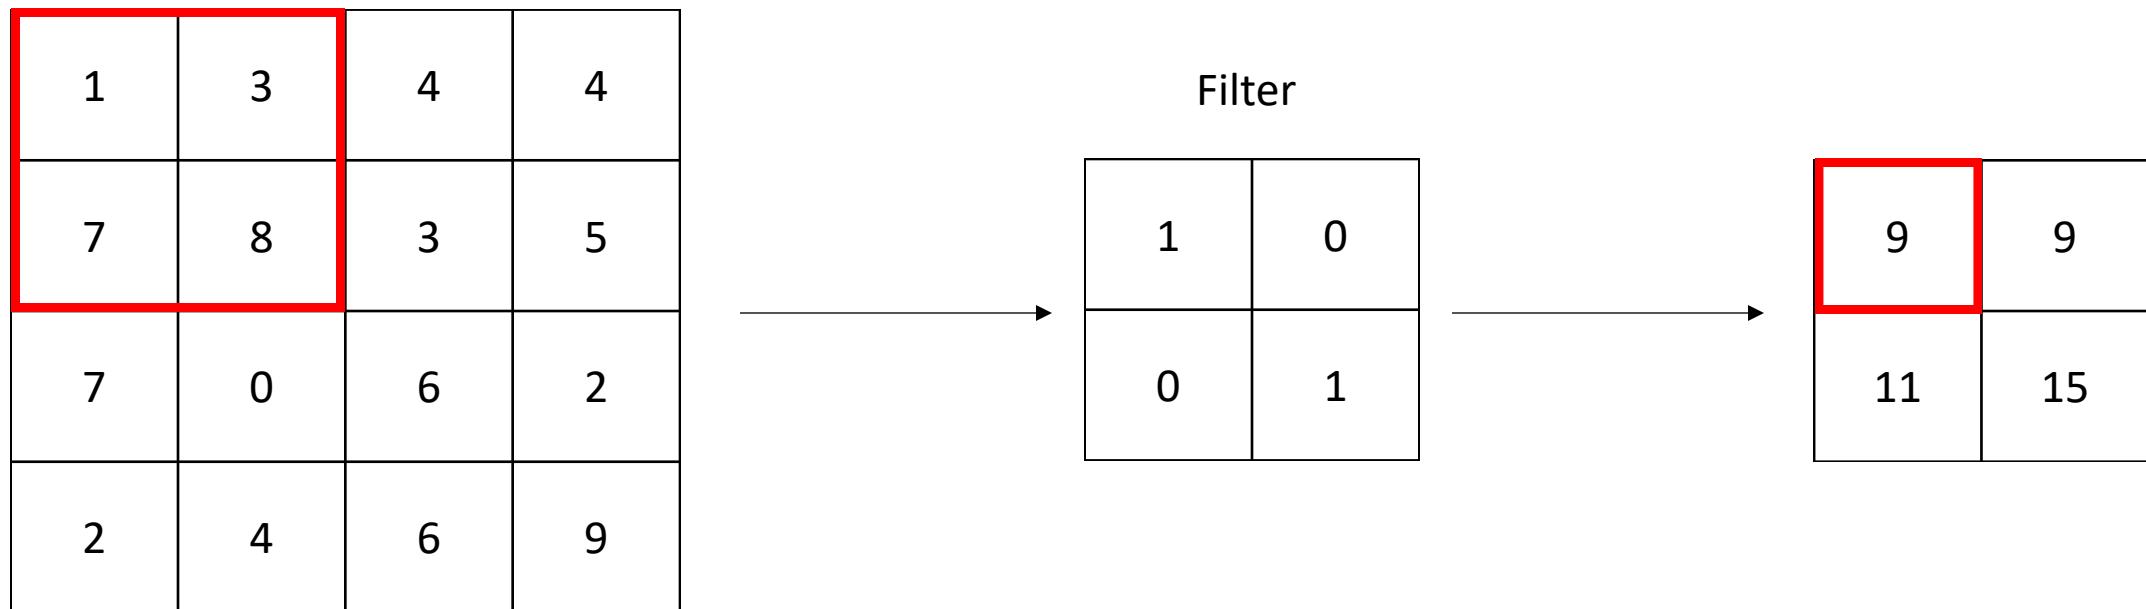
Convolutional Neural Network – Pooling Layer

1	3	4	4
7	8	3	5
7	0	6	2
2	4	6	9

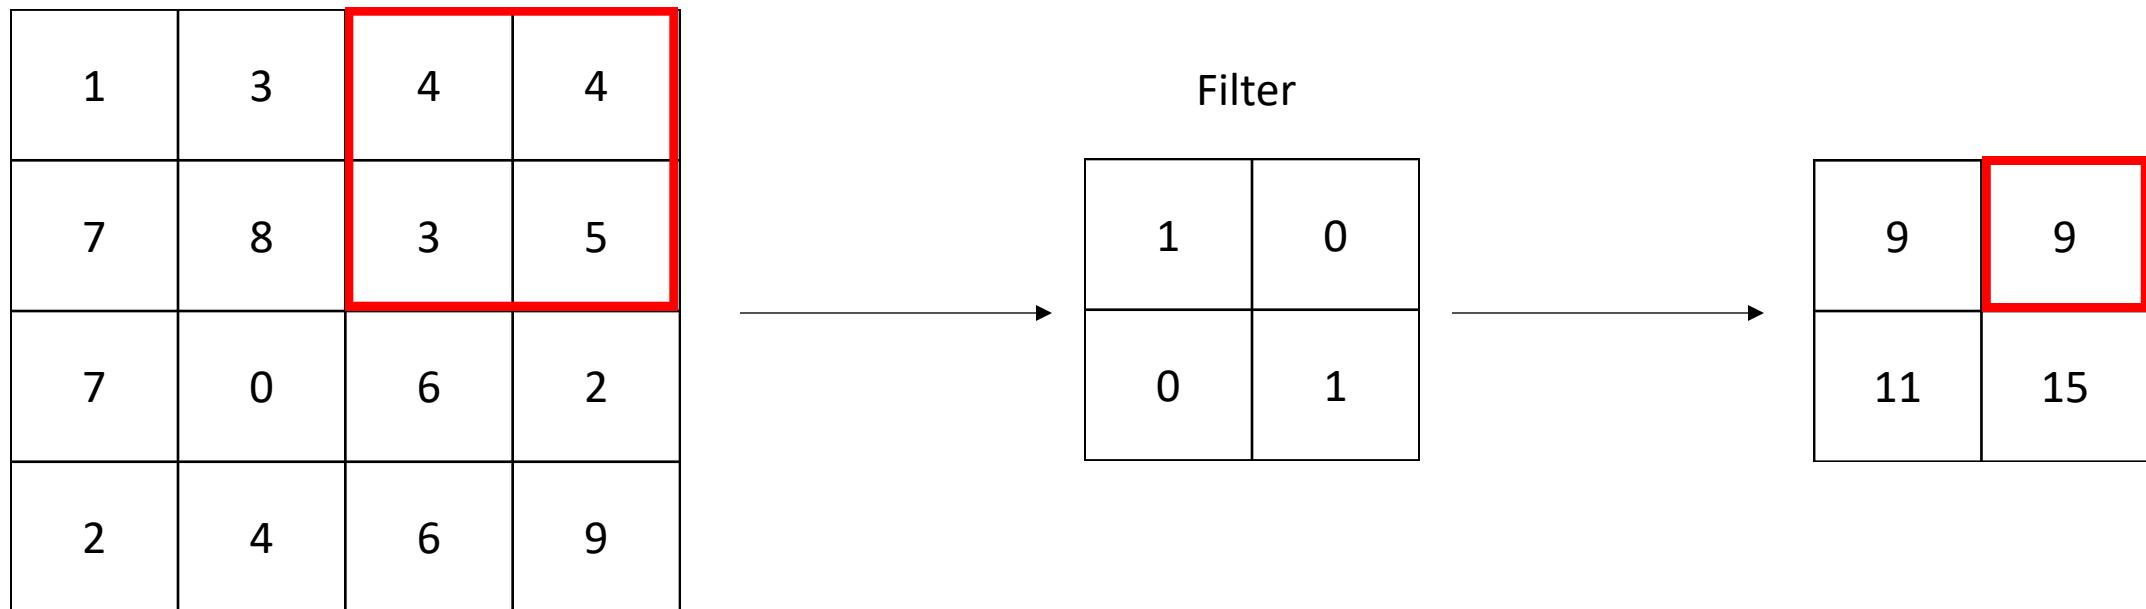


8	5
7	9

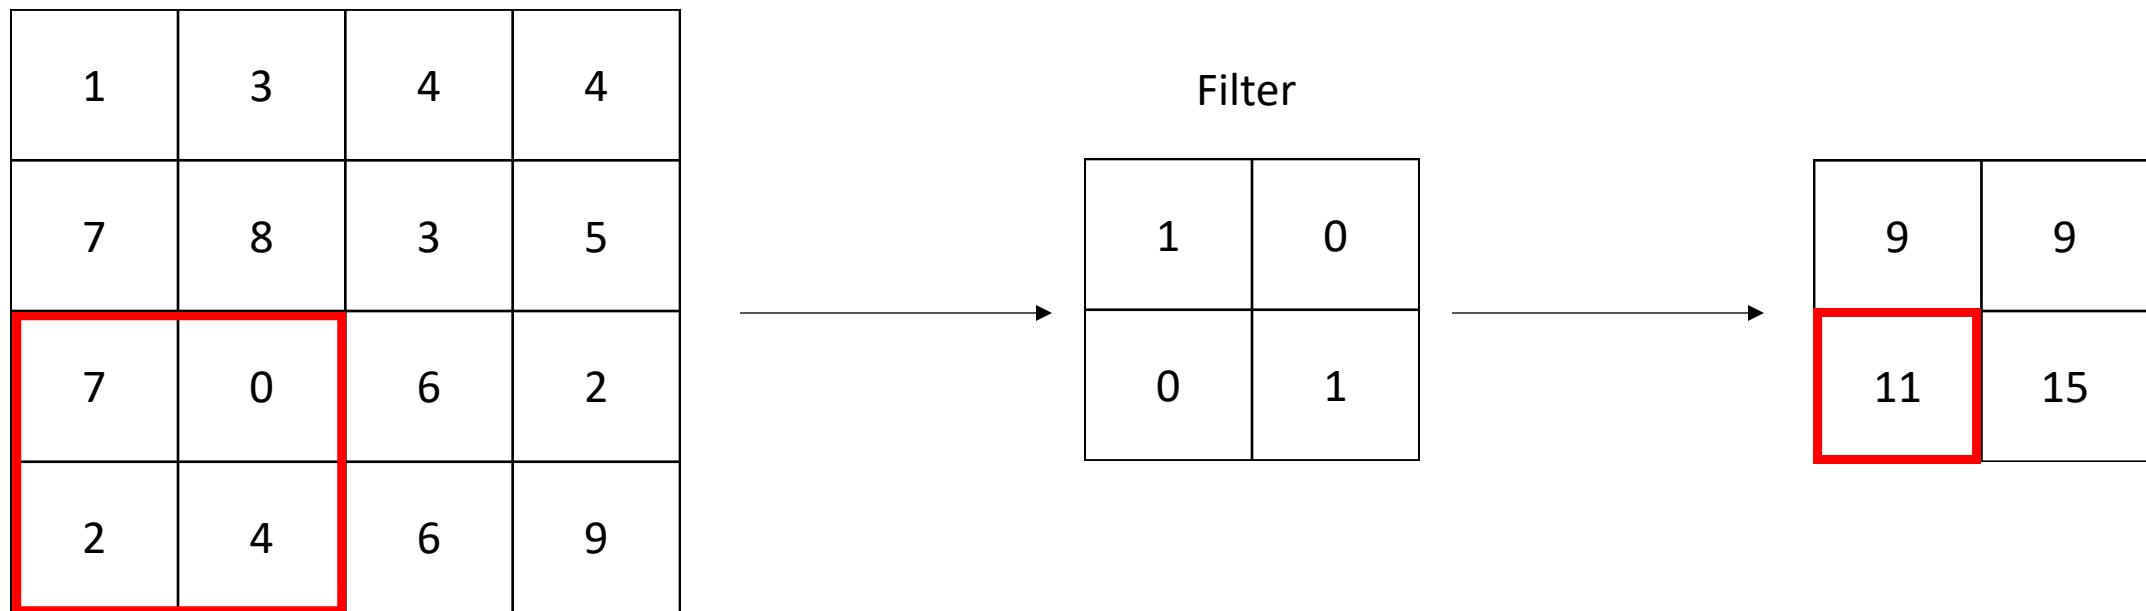
Convolutional Neural Network – Convolution Layer



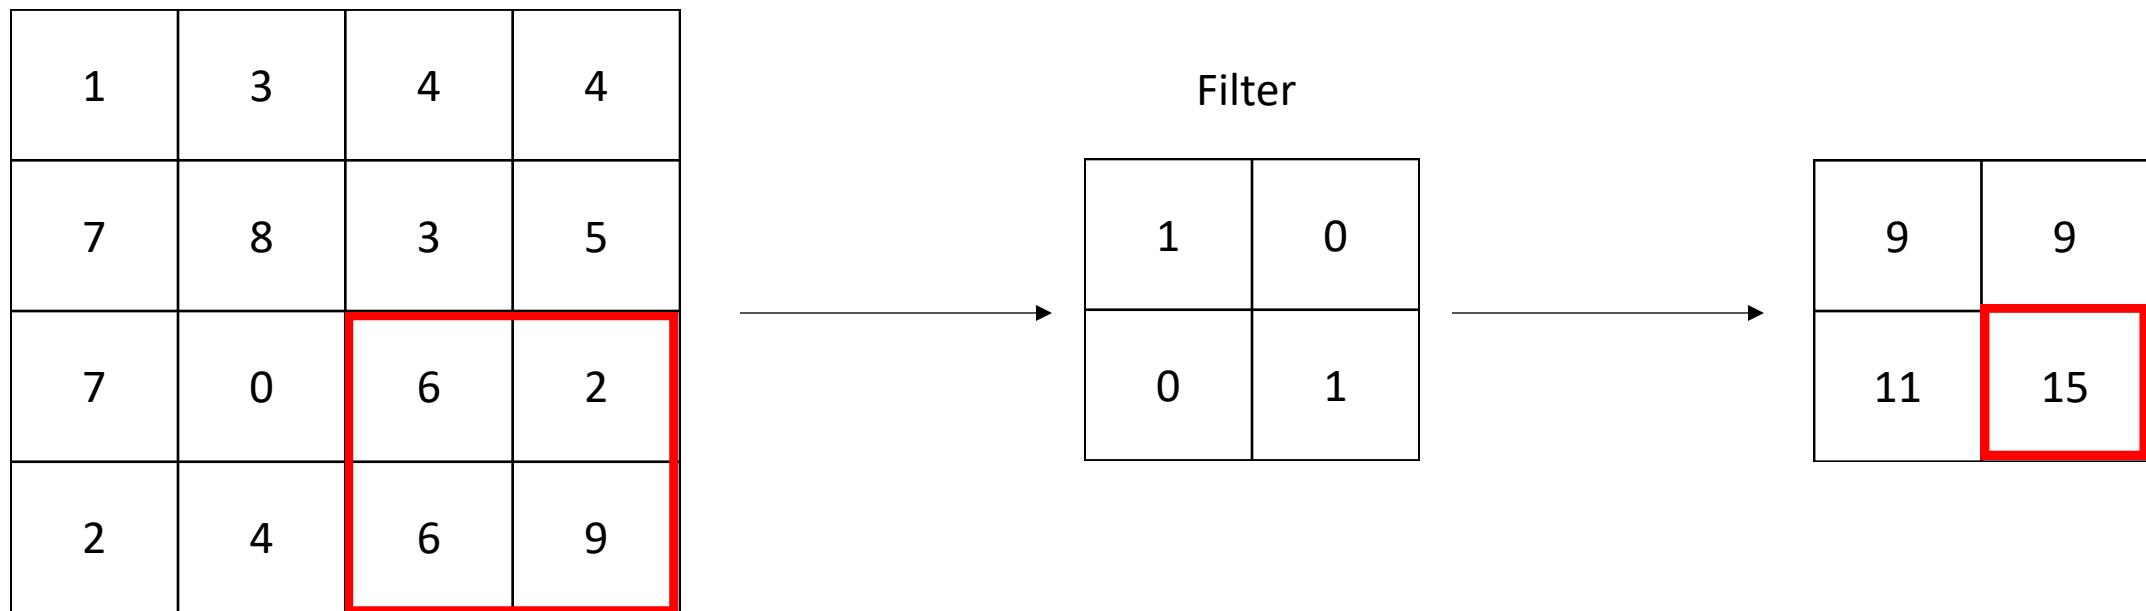
Convolutional Neural Network – Convolution Layer



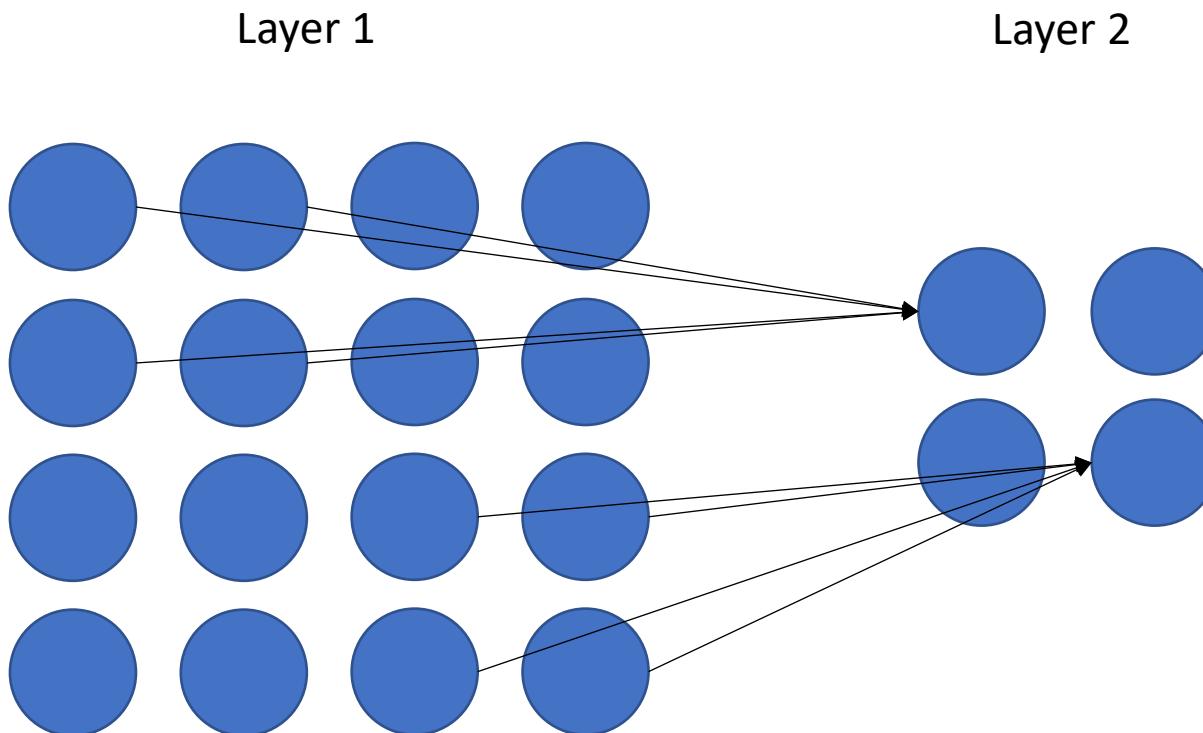
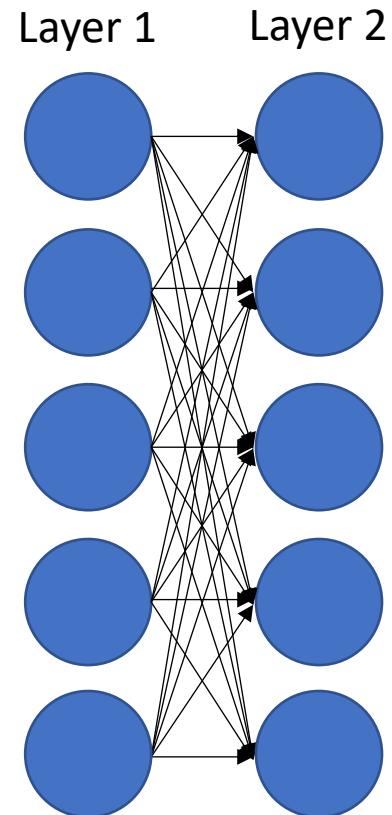
Convolutional Neural Network – Convolution Layer



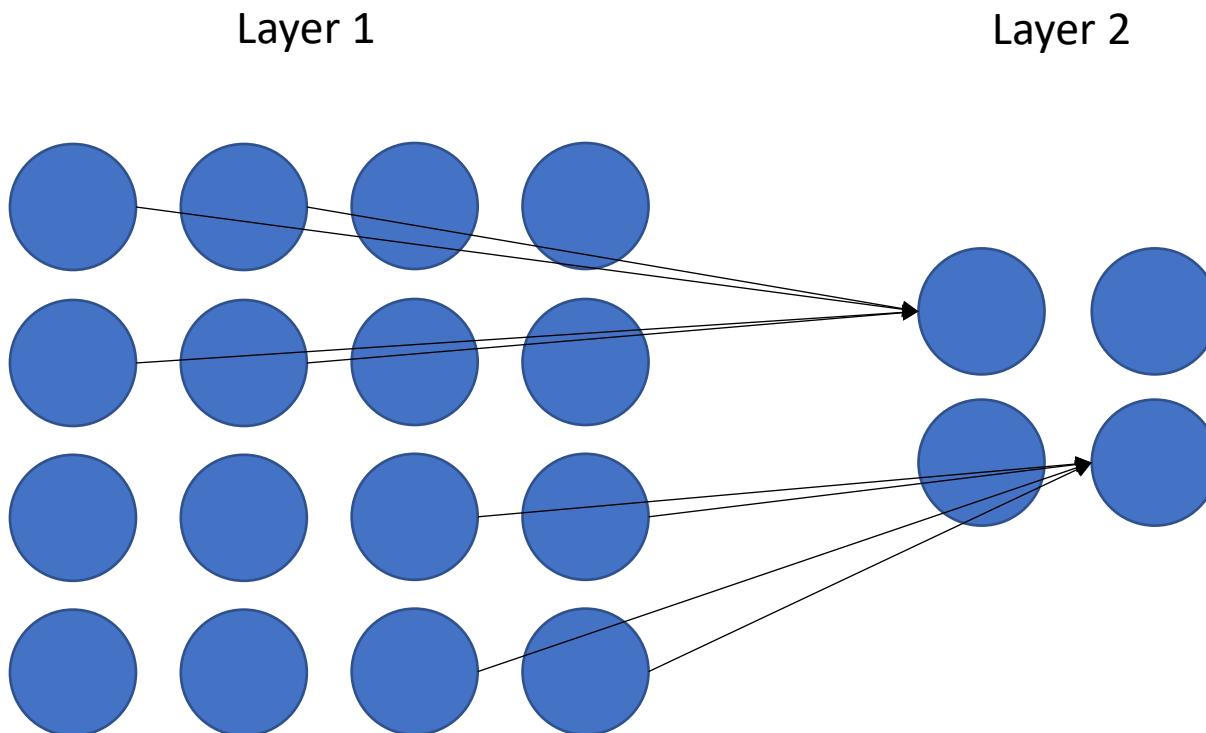
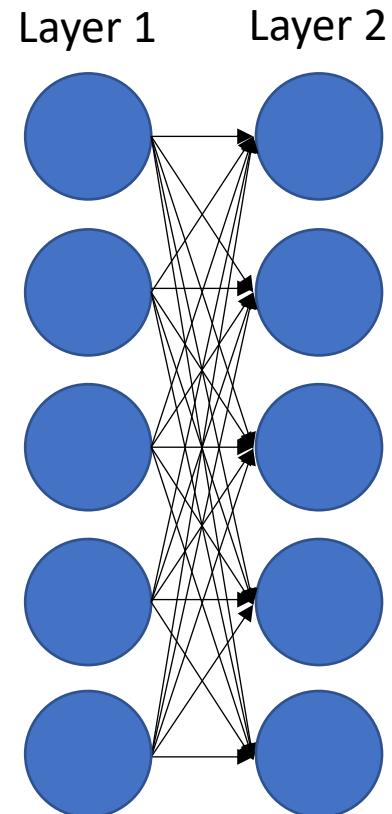
Convolutional Neural Network – Convolution Layer



Convolutional Neural Network

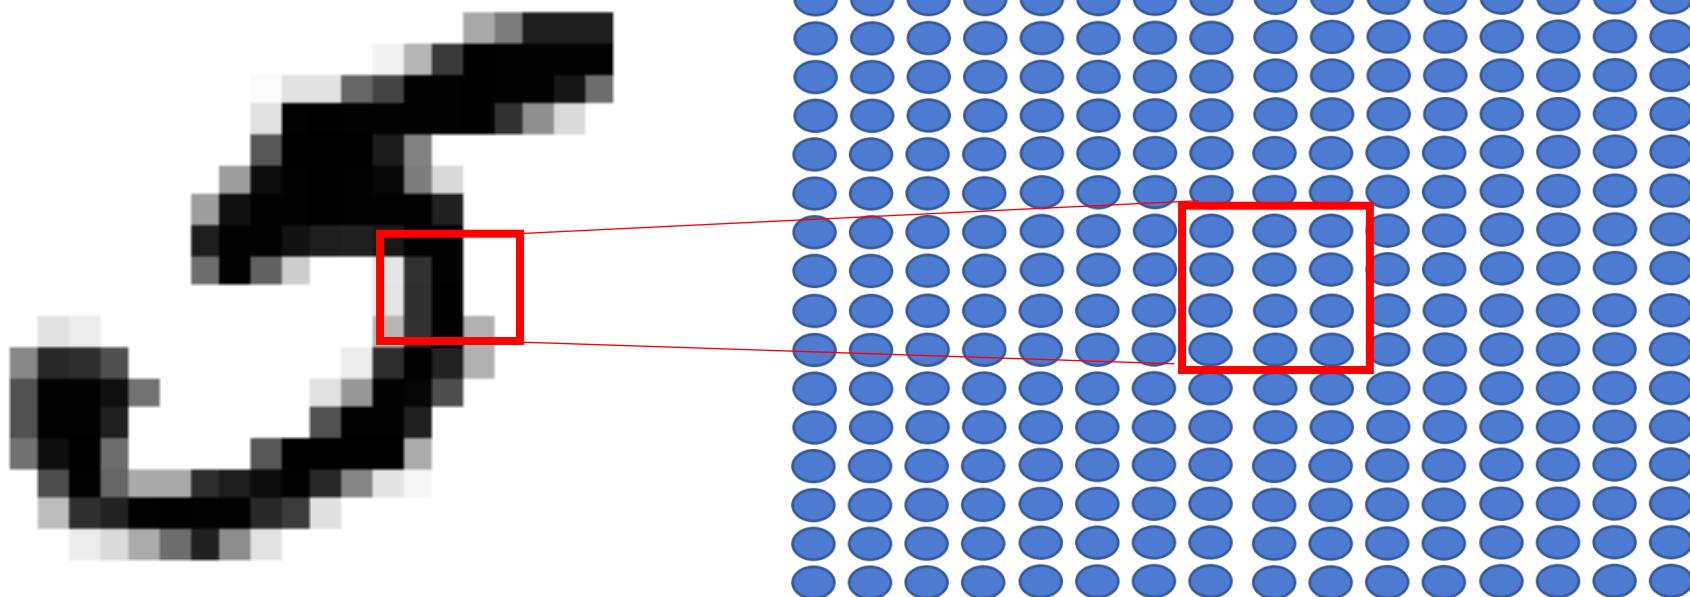


Convolutional Neural Network

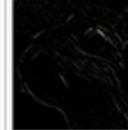
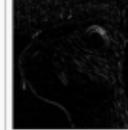


Die Gewichte entsprechen hier dem Filter und werden somit gelernt.

Convolutional Neural Network – Intuition (I)

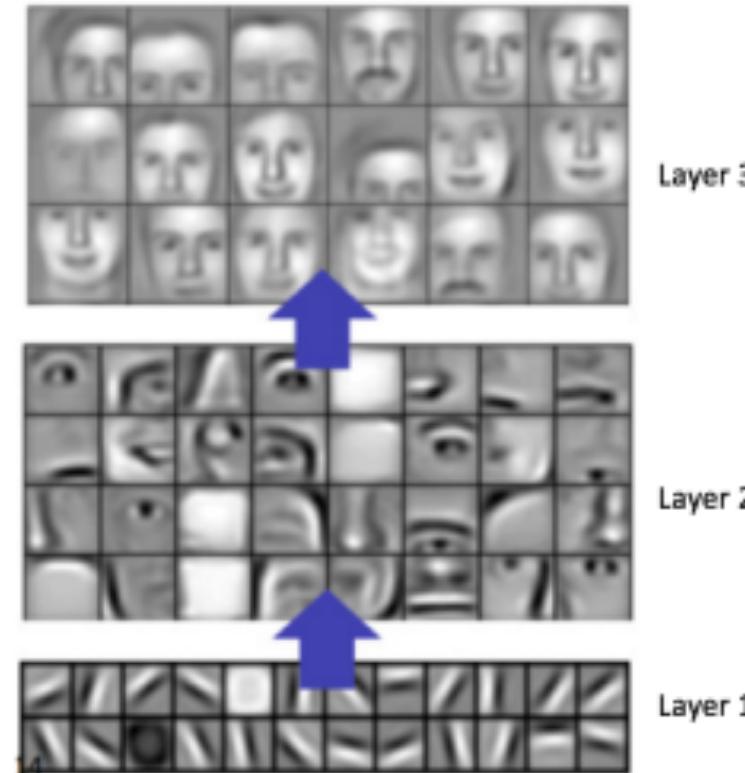


Convolutional Neural Network – Intuition (II)

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

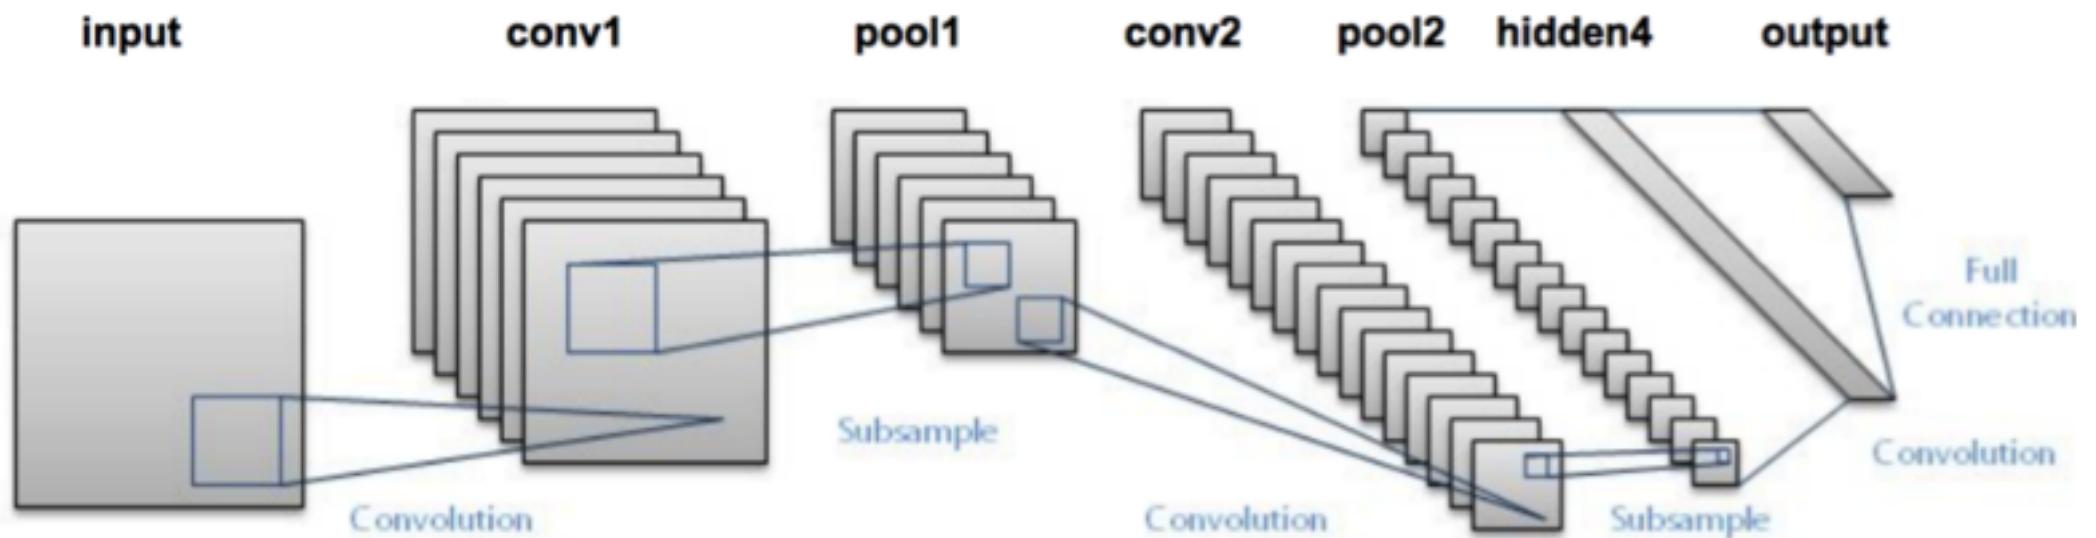
Quelle: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Convolutional Neural Network – Intuition (III)



Quelle: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Architekturen



1. Neuronen

2. Neuronale Netzwerke

3. Convolutional Neural Network (CNN)

4. Anwendungen von CNN

Bild Klassifizierung - Tensorflow Implementation – Graph

```
import tensorflow as tf

with tf.name_scope("inputs"):
    X = tf.placeholder(tf.float32, shape=[None, 784], name="X")
    X_reshaped = tf.reshape(X, shape=[-1, 28, 28, 1])
    y = tf.placeholder(tf.int32, shape=[None], name="y")

conv1 = tf.layers.conv2d(X_reshaped, filters=32, kernel_size=3,
                       strides=1, padding='SAME',
                       activation=tf.nn.relu, name="conv1")

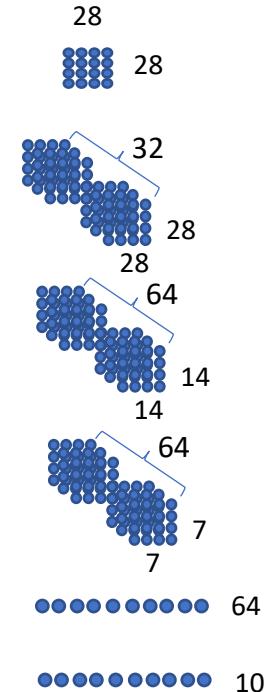
conv2 = tf.layers.conv2d(conv1, filters=64, kernel_size=3,
                       strides=2, padding='SAME',
                       activation=tf.nn.relu, name="conv2")

with tf.name_scope("pool3"):
    pool3 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="VALID")
    pool3_flat = tf.reshape(pool3, shape=[-1, 64 * 7 * 7])

with tf.name_scope("fc1"):
    fc1 = tf.layers.dense(pool3_flat, 64, activation=tf.nn.relu, name="fc1")

with tf.name_scope("output"):
    logits = tf.layers.dense(fc1, 10, name="output")

with tf.name_scope("train"):
    xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(logits=logits, labels=y)
    loss = tf.reduce_mean(xentropy)
    optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)
    training_op = optimizer.minimize(loss)
```



y_0, y_1, \dots, y_9

Bild Klassifizierung - Tensorflow Implementation - Session

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("/tmp/data/")

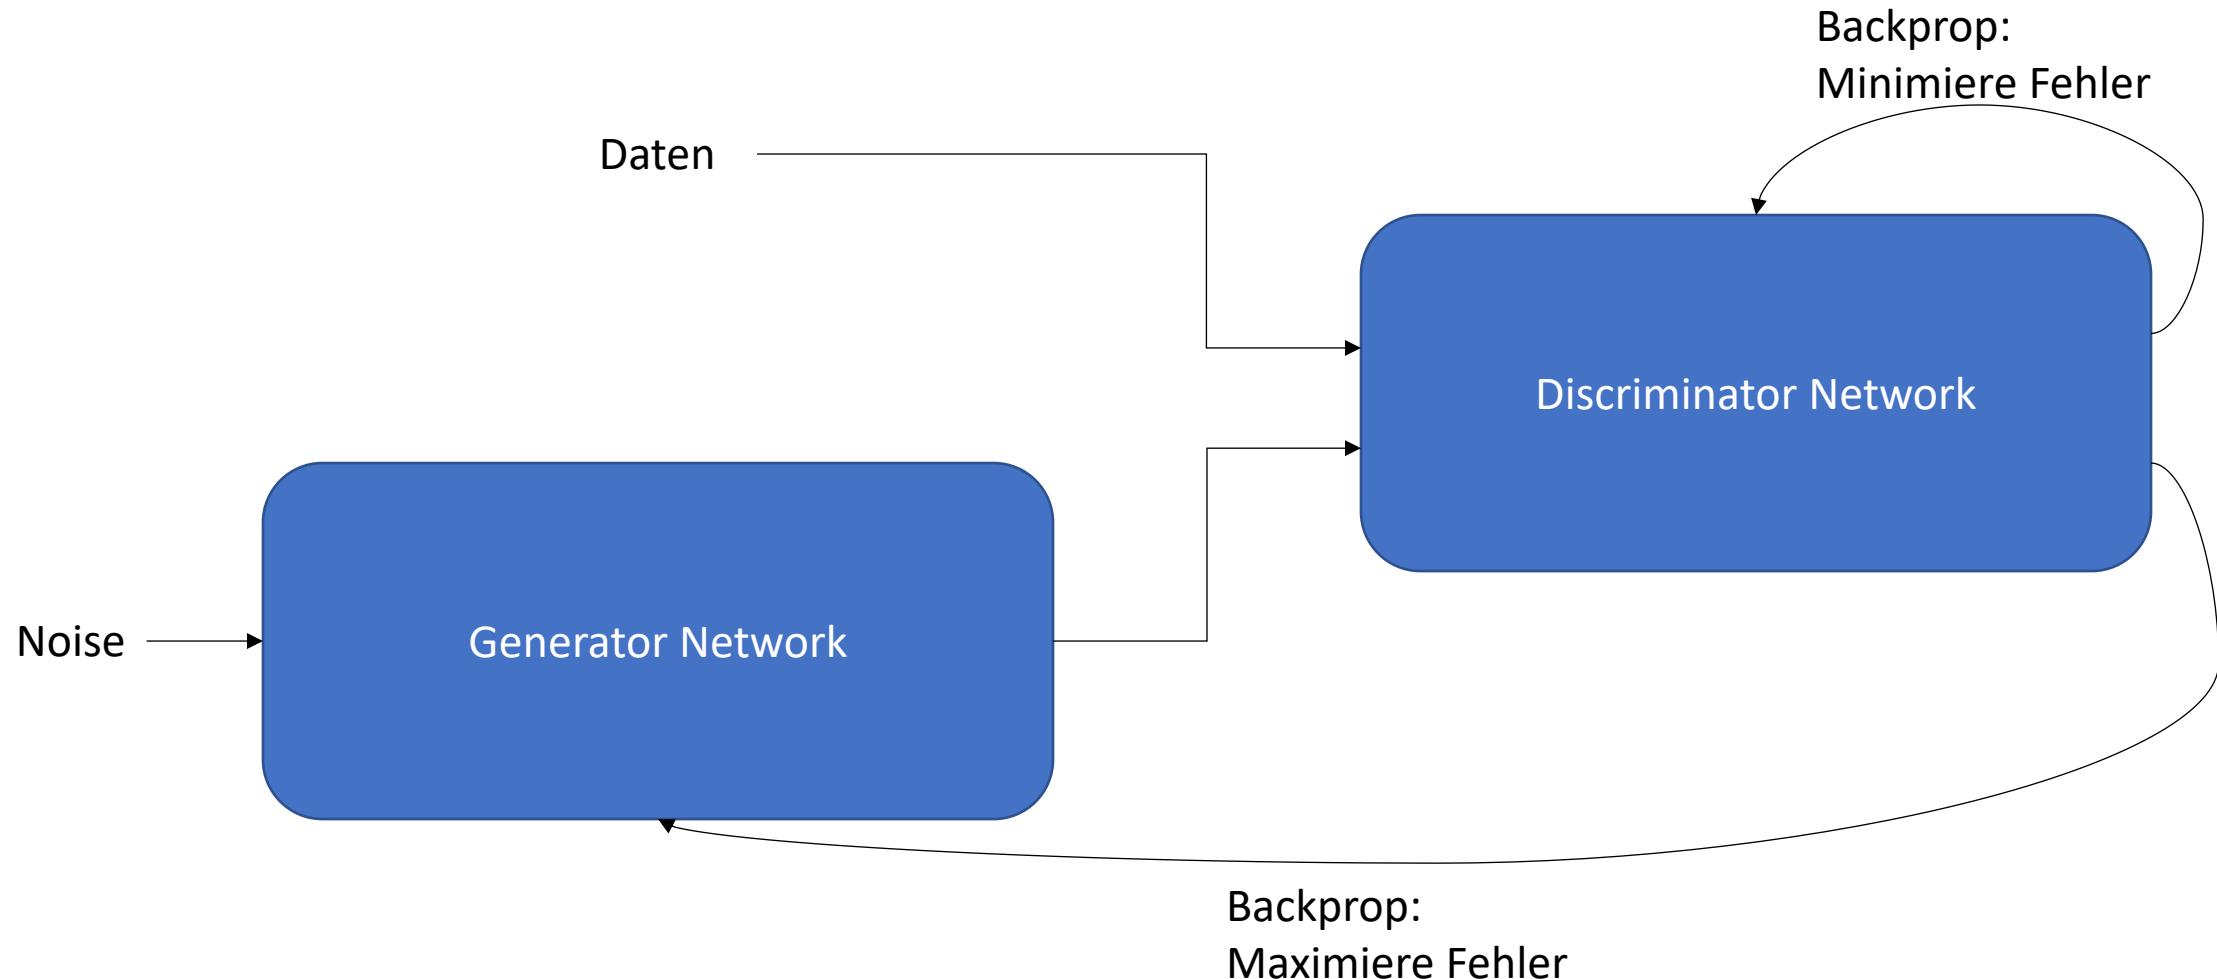
init = tf.global_variables_initializer()

n_epochs = 10
batch_size = 100

with tf.Session() as sess:
    init.run()
    for epoch in range(n_epochs):
        for iteration in range(mnist.train.num_examples // batch_size):
            X_batch, y_batch = mnist.train.next_batch(batch_size)
            sess.run(training_op, feed_dict={X: X_batch, y: y_batch})
            acc_train = accuracy.eval(feed_dict={X: X_batch, y: y_batch})
            acc_test = accuracy.eval(feed_dict={X: mnist.test.images, y: mnist.test.labels})
            print(epoch, "Train accuracy:", acc_train, "Test accuracy:", acc_test)

0 Train accuracy: 0.97 Test accuracy: 0.9789
1 Train accuracy: 0.99 Test accuracy: 0.9844
2 Train accuracy: 0.99 Test accuracy: 0.9867
3 Train accuracy: 0.99 Test accuracy: 0.9872
4 Train accuracy: 0.99 Test accuracy: 0.9879
5 Train accuracy: 0.99 Test accuracy: 0.9899
6 Train accuracy: 0.99 Test accuracy: 0.9885
7 Train accuracy: 1.0 Test accuracy: 0.9876
8 Train accuracy: 1.0 Test accuracy: 0.9891
9 Train accuracy: 1.0 Test accuracy: 0.9834
```

Generative Adversarial Networks

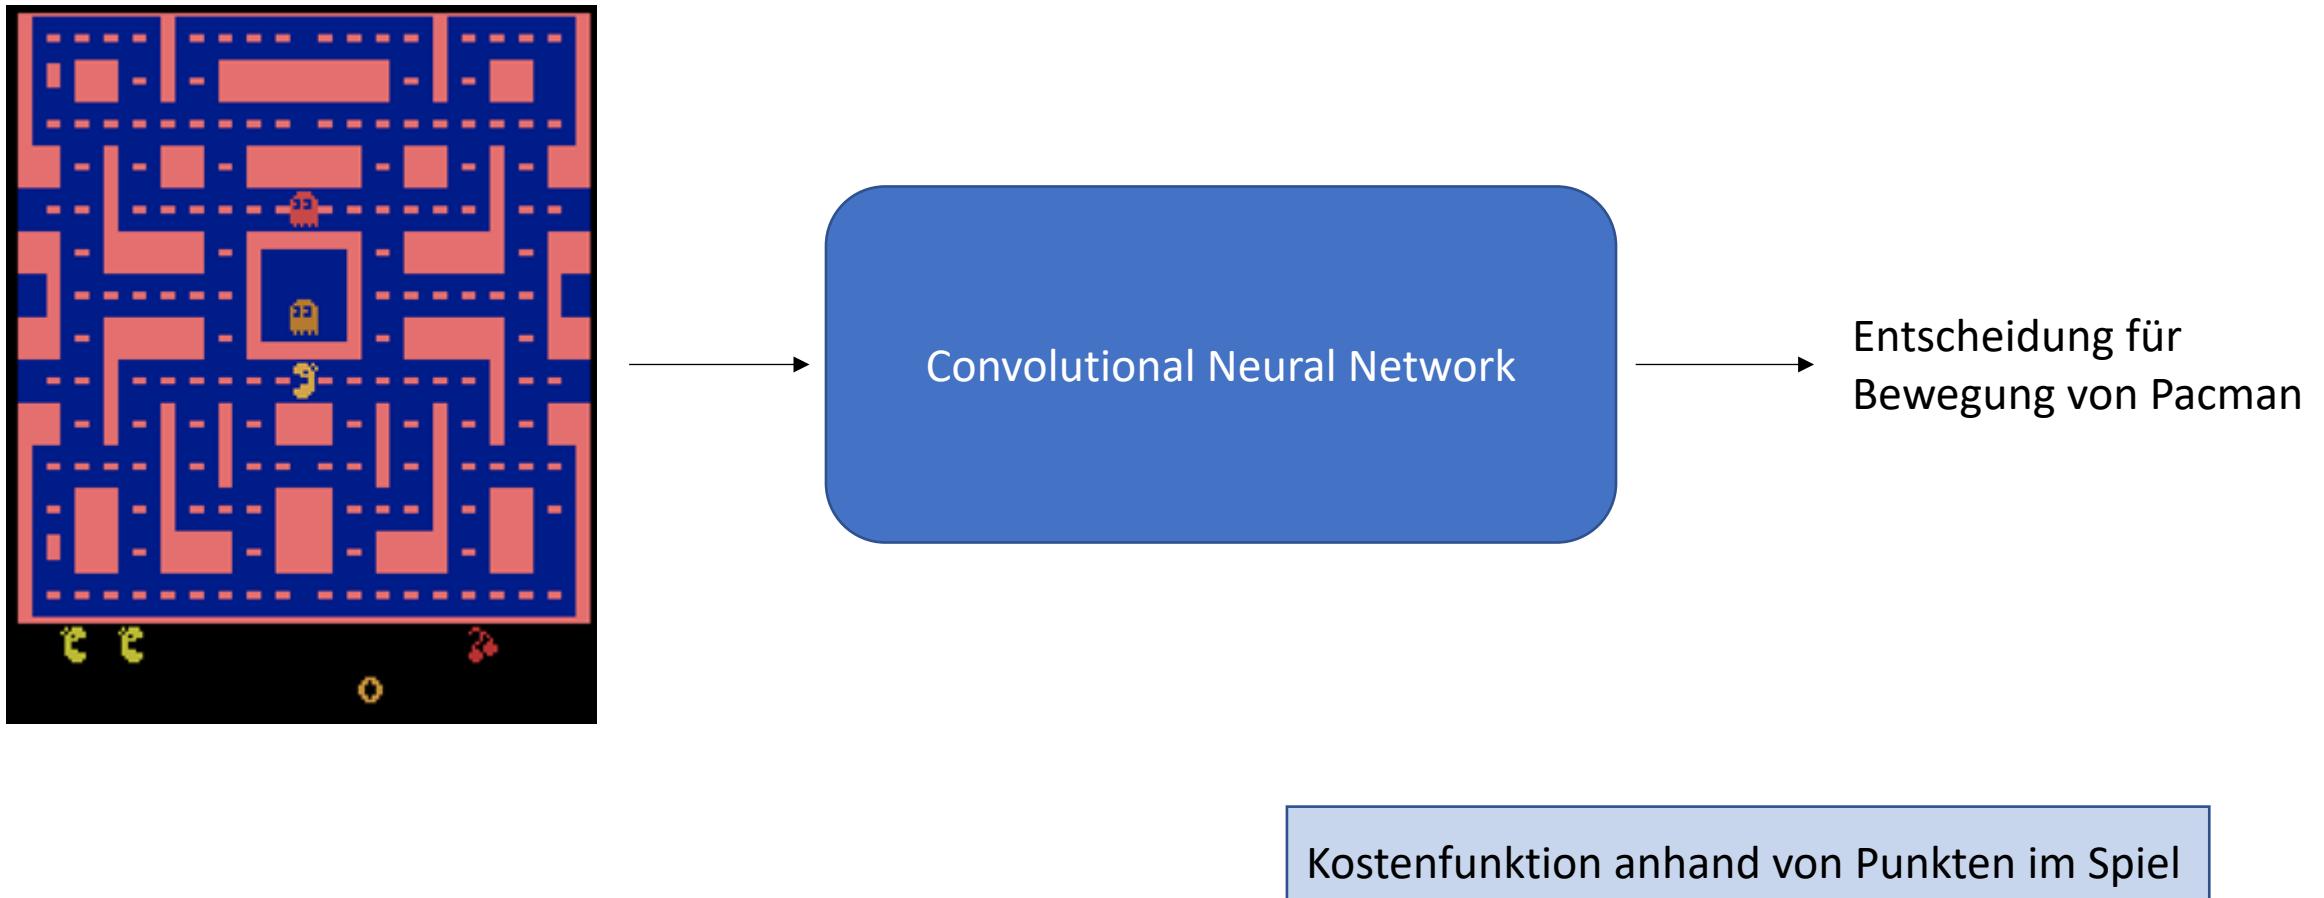


Generative Adversarial Networks



Quelle: <https://towardsdatascience.com/mangagan-8362f06b9625>

Reinforcement Learning



What else?

- Aktivierungsfunktion
- Batch Normalization
- Dropout
- Vortraining
- Optimierungsalgorithmen
- ...

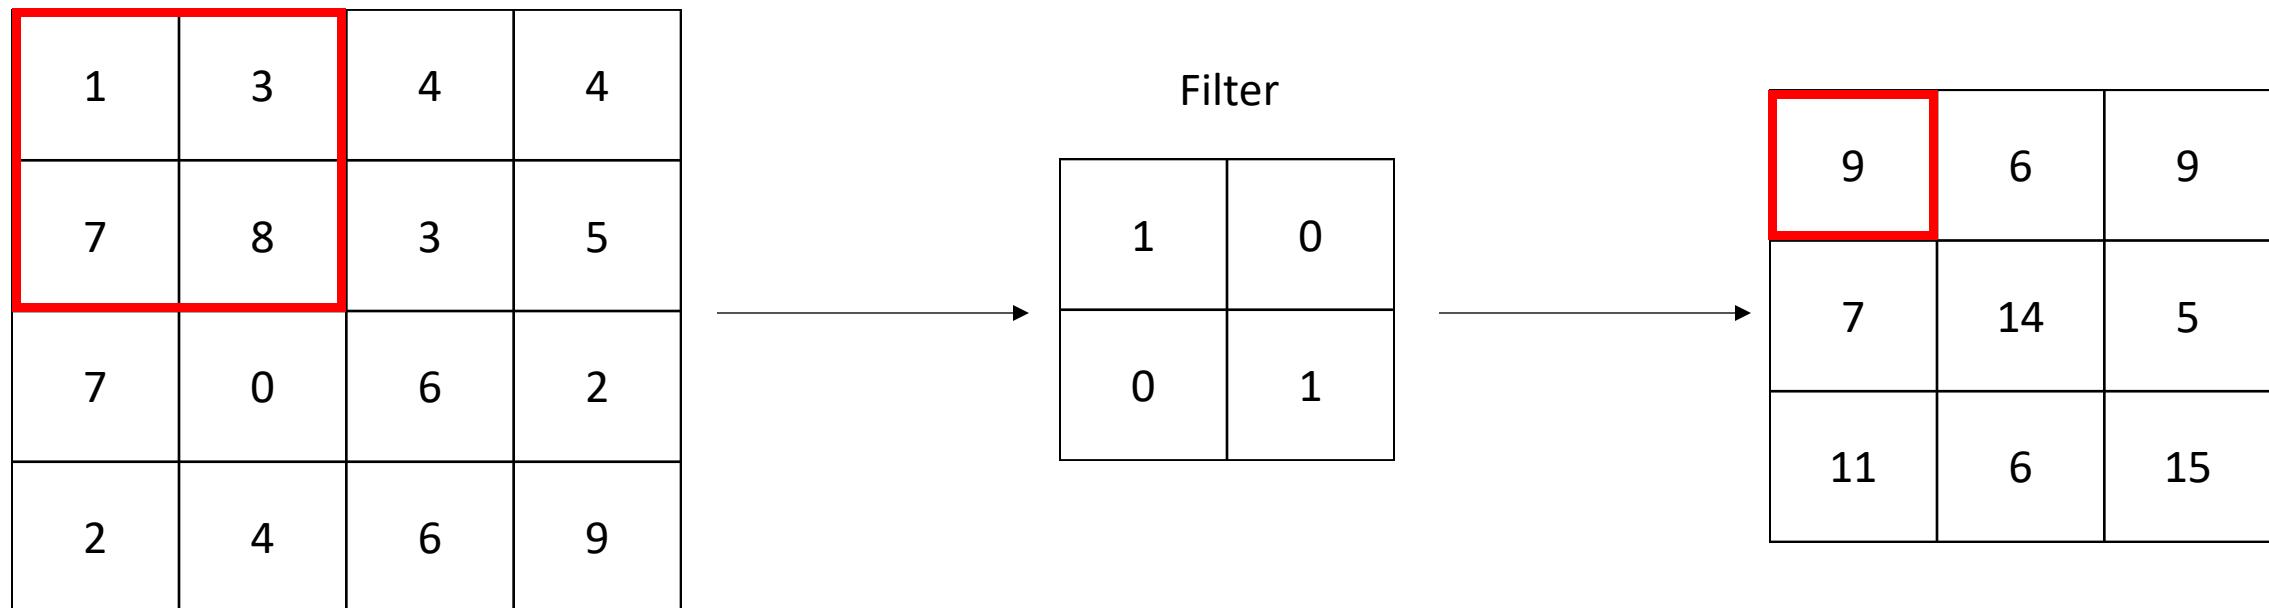


Vielen Dank für die Aufmerksamkeit

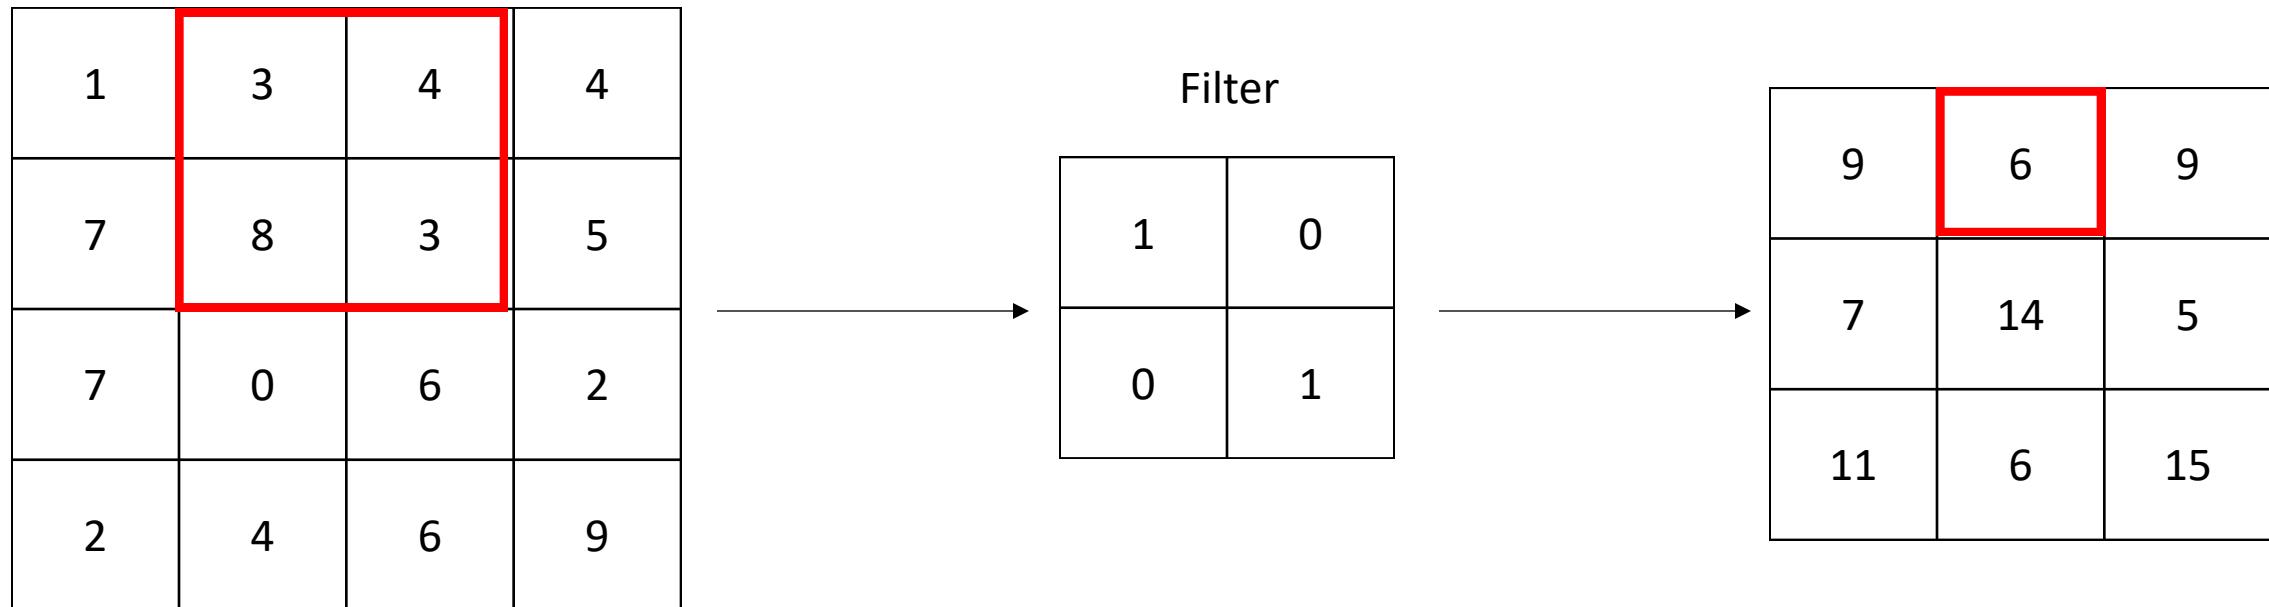
Nicolas Kuhaupt
Nicolaskuhaupt@aol.com
Data Scientist Fraunhofer IEE

Backup

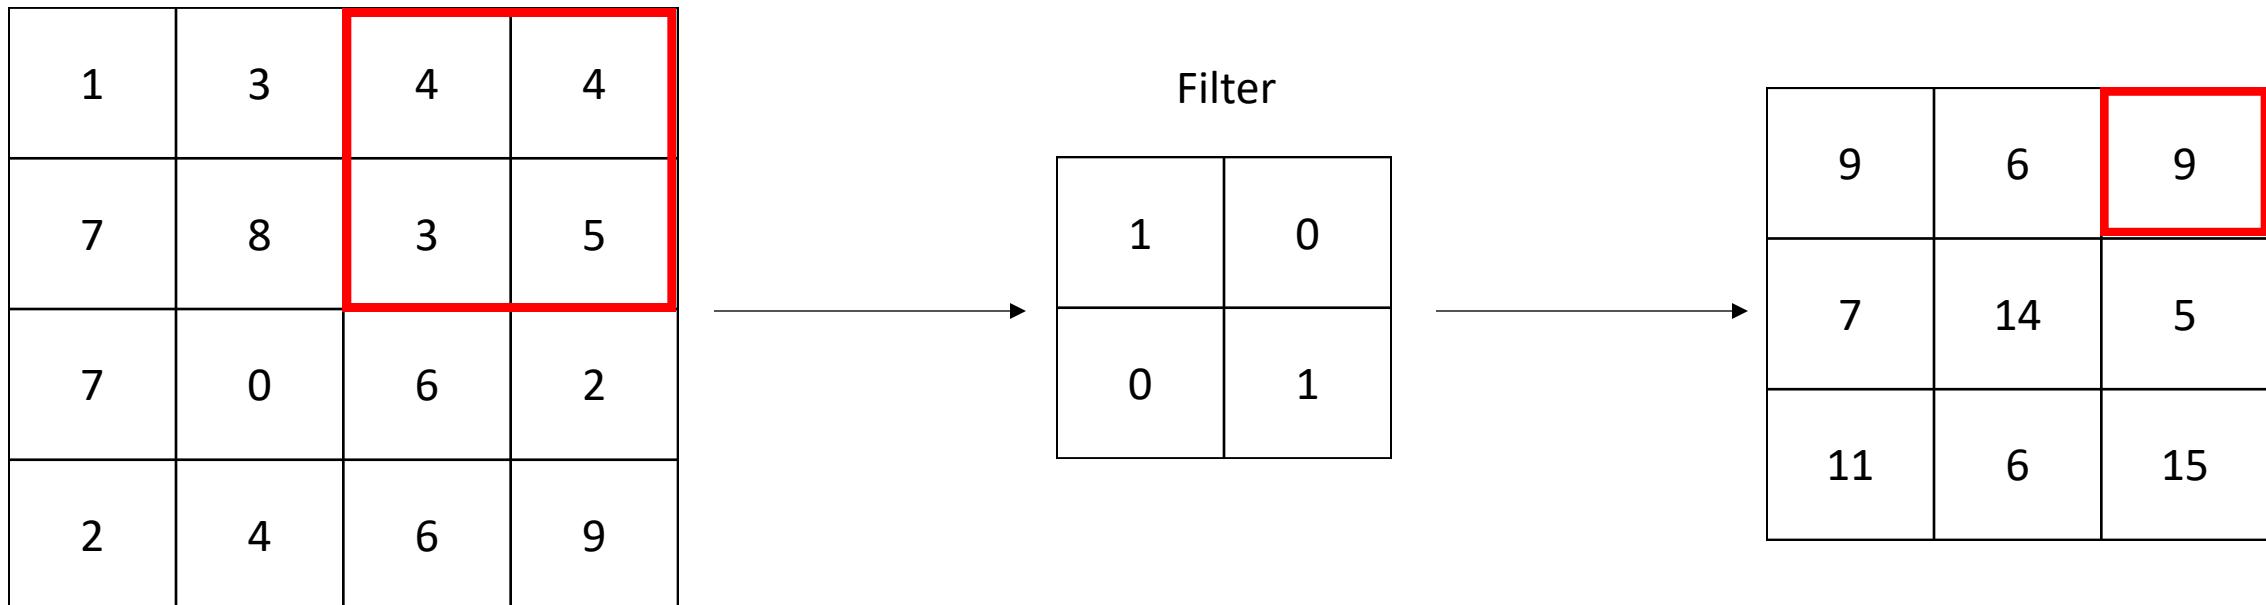
Convolutional Neural Network – Stride



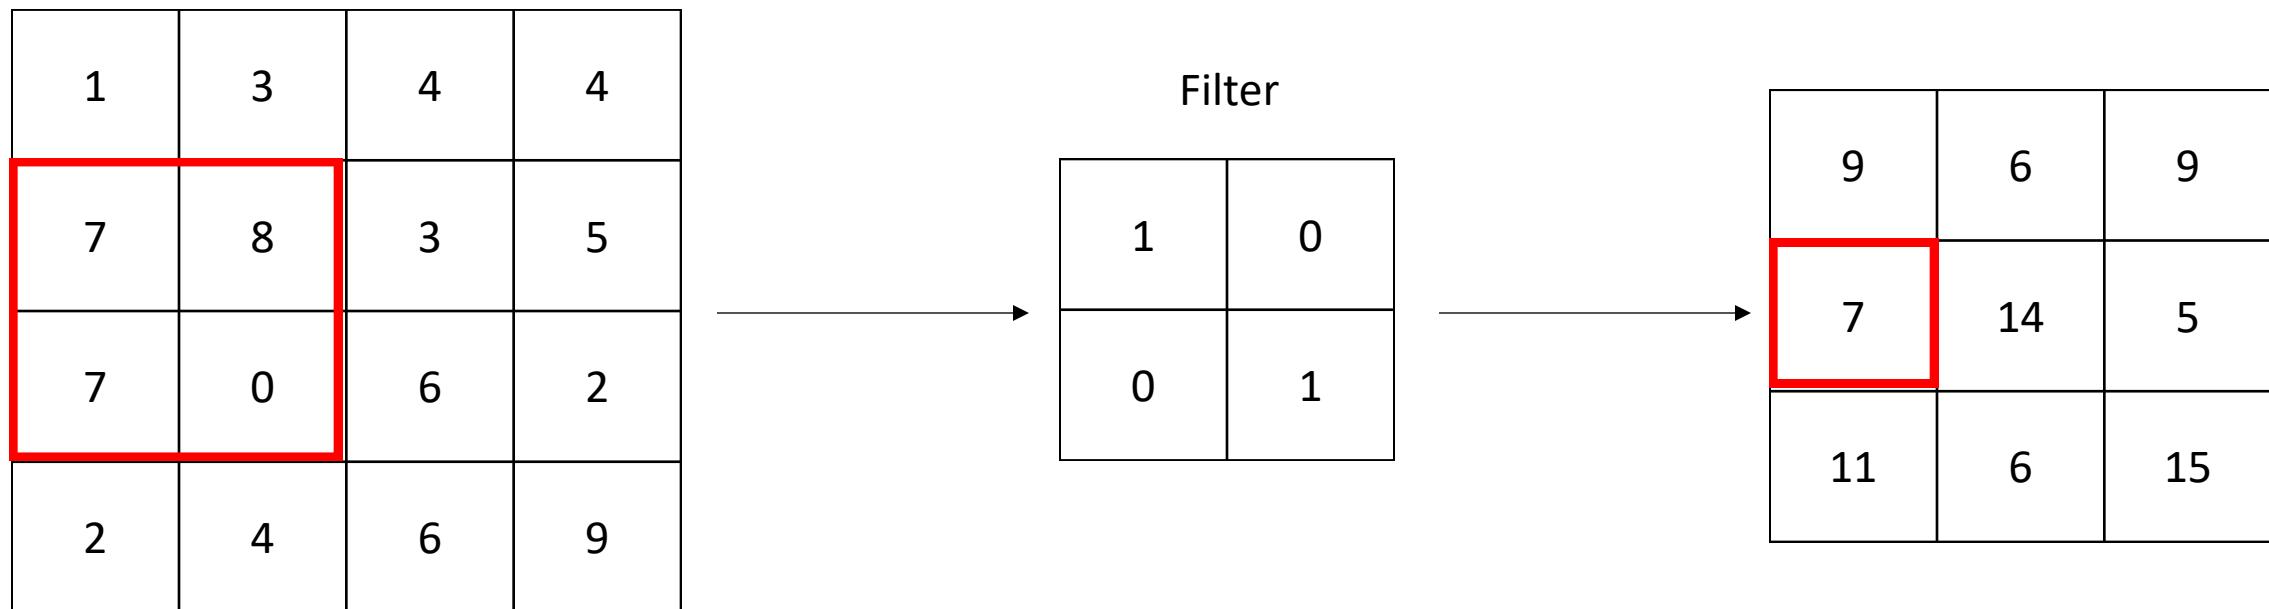
Convolutional Neural Network – Stride



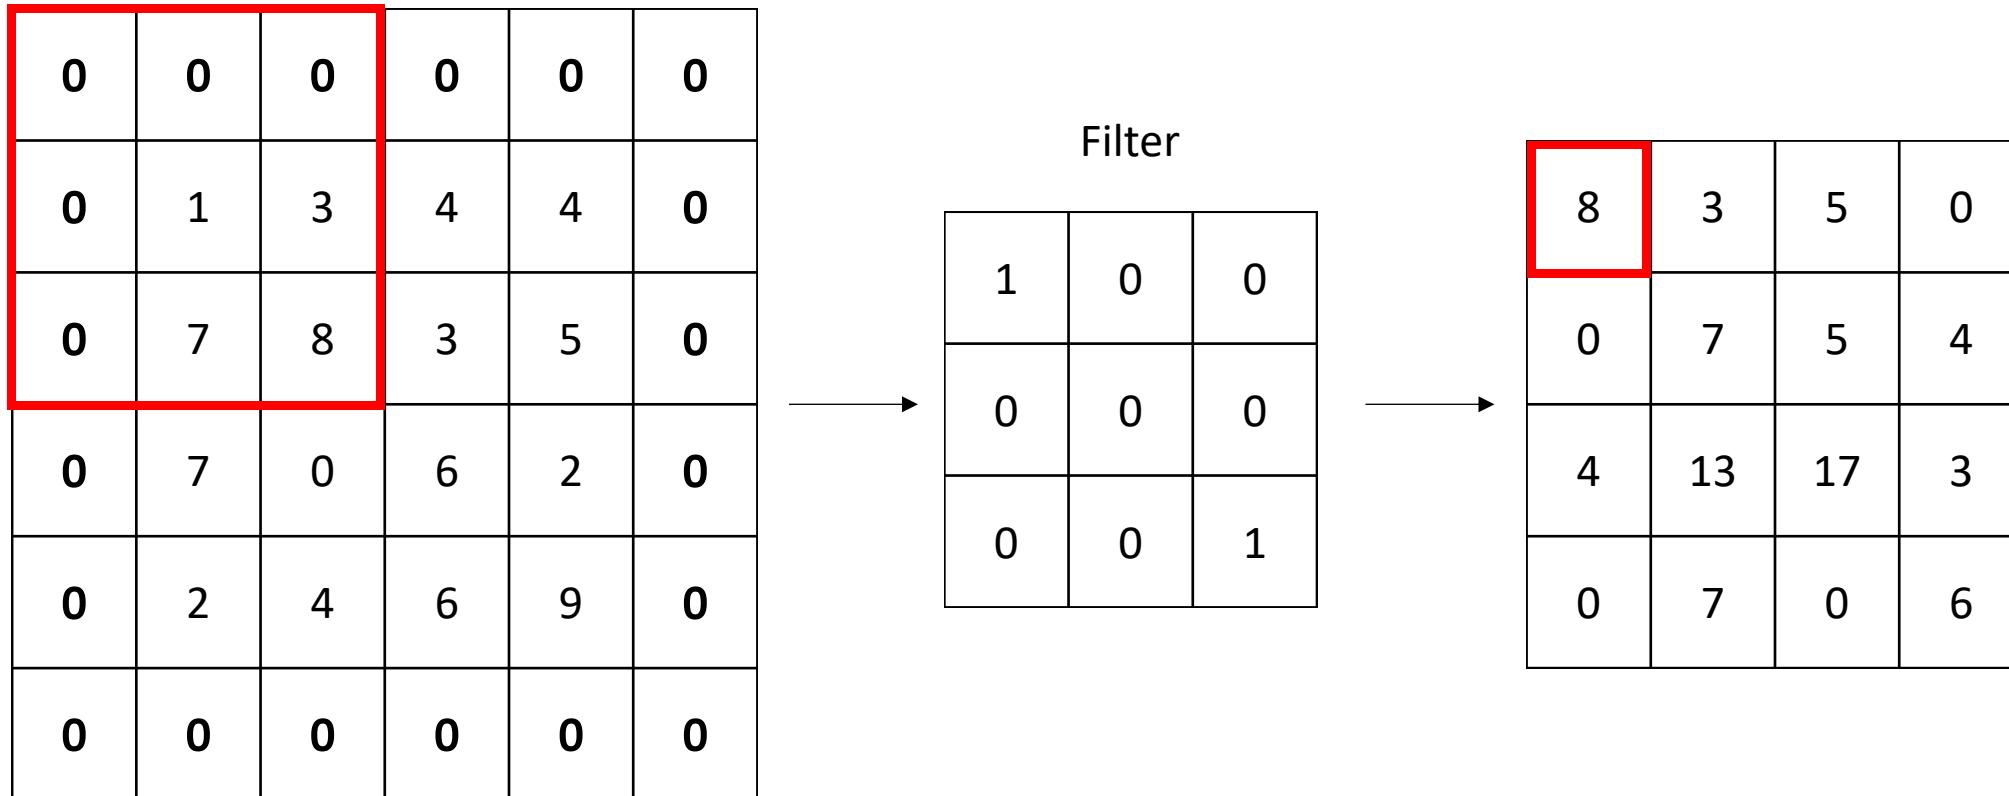
Convolutional Neural Network – Stride



Convolutional Neural Network – Stride



Convolutional Neural Network – Padding



Convolutional Neural Network – Padding

